

# Roadmap

## Instruction Level Parallelism (ILP)

1. Review of Simple Pipelined Processors
2. Dynamically Scheduled Pipelines
3. Superscalar and Multithreaded Processors

### 1. Simple Pipelining

- Baseline design:
  - Introduce a simple RISC Instruction Set Architecture (ISA) (called DLX)
  - Non-pipelined DLX implementation
- 5-Stage Pipelined DLX Processor
- Hazards and their solutions

## Roadmap

### Instruction Level Parallelsim (ILP)

Overlapped execution of instructions from a sequential program

- Programmer obliviousness
- Single-threaded applications
- Compiler support

### Limits

ILP Wall

Memory Wall

Power Wall

# Simplified DLX Instruction Set

DLX: Idealized RISC processor (similar to MIPS, ARM)

- Load/Store architecture using base + offset addressing
- 32 bit word size aligned at word address boundaries
- 32-bit memory addresses (aligned)

## Registers

- 32 32-bit Integer GPRs R0 .... R31
- 32 32-bit Floating Point Registers F0, .. F31
- No condition codes


## Instructions

- Integer ALU
- Floating Point (FP)
- Memory Access
- Program Control :


# Instruction Formats

All instructions are one word long (32 bit aligned)


**R-R Instructions:** `add rd, rs, rt`



**R-I Instructions:** `addi rt, rs, d`



**Load/Store/Branch Instructions:** `lw rt, d(rs), sw rt, d(rs), beq rs, rt, d(PC)`



# Simplified Integer DLX Instruction Set

## 1. ALU instructions: (ADD, SUB, AND, OR, XOR, ....)

**RR mode:**      ADD      rd, rs, rt                      |    $[rd] = [rs] + [rt]$

**Example:**              ADD      R2, R4, R5

The 32-bit contents of registers **rs** and **rt** are added, and the sum is written to register **rd**.

**RI mode:**      ADDI      rt, rs, d                      |    $[rt] = [rs] + \text{EXT}(d)$

**Example:**              ADDI      R2, R4, 1000

The 16-bit value **d** is sign-extended to 32 bits and added to the 32-bit contents of register **rs**. The sum is written to register **rt**.

# Simplified Integer DLX Instruction Set

## 2. Memory reference instructions: Load (LW) and Store (SW)

LW $rt, d(rs)$	$ea = EXT(d) + (rs); (rt) = MEM[ea]$
SW $rt, d(rs)$	$ea = EXT(d) + (rs); MEM[ea] = (rt)$

### Example:

LW $R5, 0(R2)$	$R5$ gets the word whose <i>address</i> is stored in $R2$
SW $R6, 1000(R2)$	$R5$ written to memory location with address   $1000 + \text{value in } R2$

- **LW:** Reads a word from memory at effective address  $[ea]$  and writes it to register  $rt$
- **SW:** Writes the contents of register  $rt$  to memory at effective address  $[ea]$
- **Effective Address  $ea$** 
  - 16-bit displacement  $d$  is sign-extended to 32 bits and added to the contents of base register  $rs$  to get the effective address.

# Simplified Integer DLX Instruction Set

## 3. Control instructions

### Conditional Branch

if **condition** is TRUE  
    go to Target Address;  
else  
    continue with next in-line instruction

- **Compare values in two registers** (e.g. BEQ, BNE, BGT, BLT, BGW, BLE, ... )
- **Compare with 0** (e.g. BEQZ, BNEZ, BGTZ, BLTZ, BGEZ, BLEZ, ... )

### Example

BEQ   rs, rt, d(PC)   | Go to Target Address if contents of registers rs and rt are equal

BEQZ   rs, d(PC)       | Go to Target Address if contents of register rs equals zero

**Target Address** = PC + 4 + Extended(d)

16-bit displacement d is sign-extended to 32 bits

PC is address of the branch instruction

Blank Slide



# DLX Implementation

## Functional Modules

- Register File (REG)
- Instruction Memory (IM)
- Data memory (DM)
- Arithmetic Logic Unit (ALU)
- Decoder, Program Counter

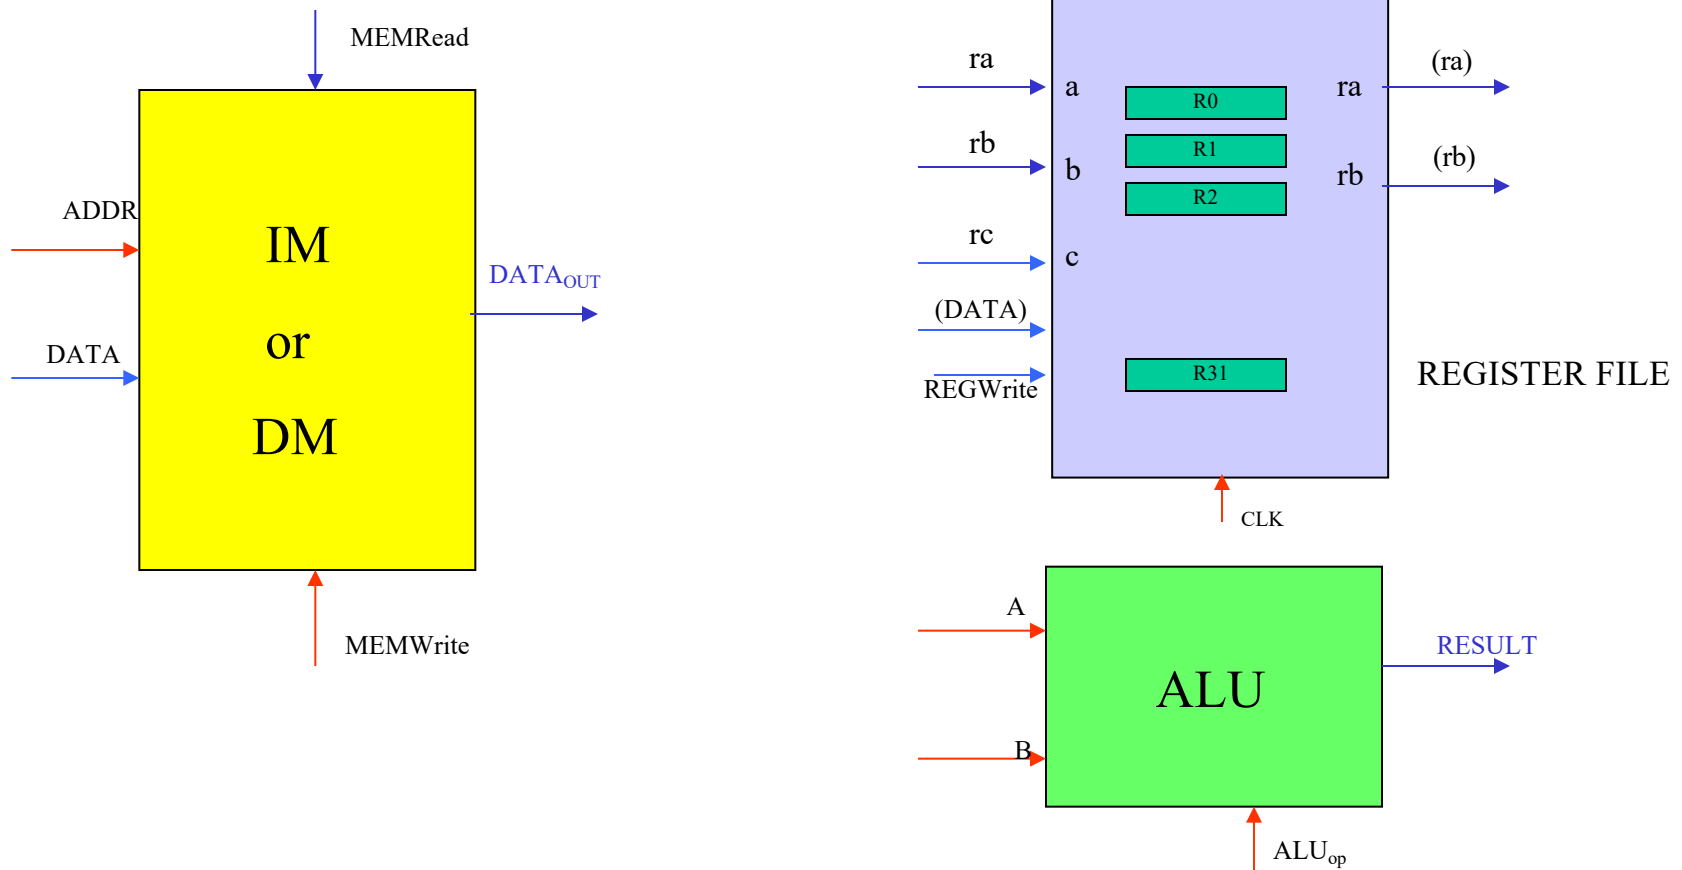
## DataPaths

Path of information in executing an instruction: sequence of FMs

## Timing and Control

Enforce sequencing in the datapath

## Functional Units



**Decoder:** Combinatorial Circuit that generates control signals from instruction

**PC:** Register holds address of current instruction. Changes value at next clock edge

## Hypothetical Single-cycle Implementation of DLX

Each instruction completes in 1 clock cycle

- Registers (PC + Register File) have stable values after clock edge

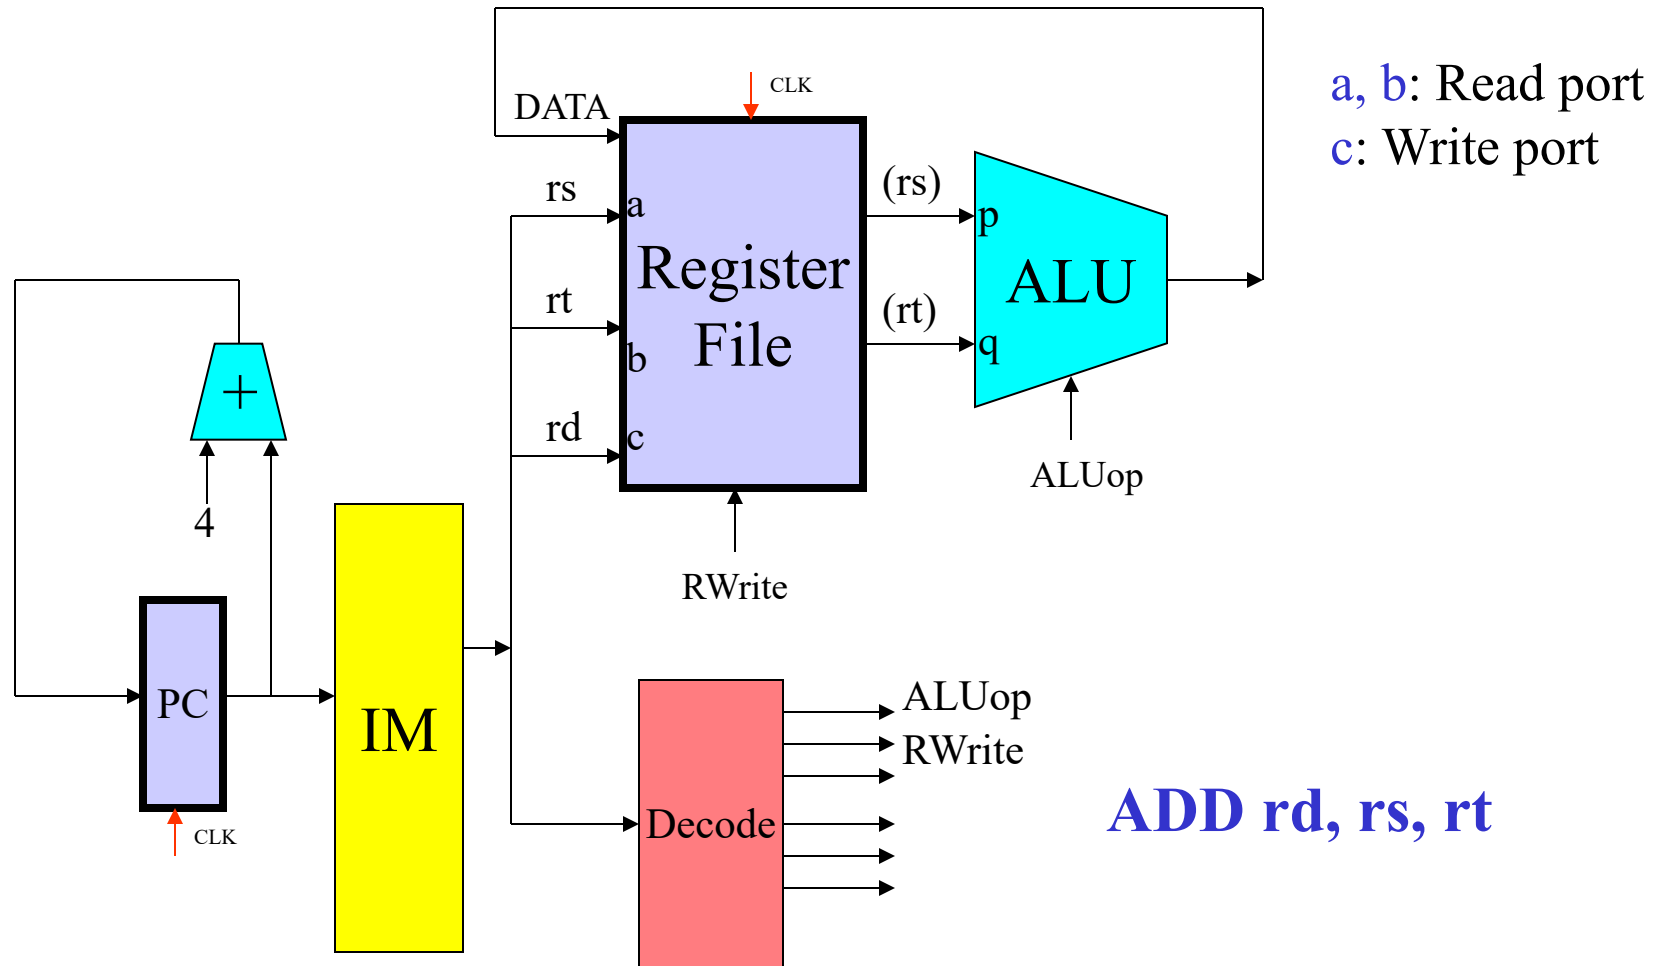
During clock cycle:

1. Read Instruction from Instruction memory (IM)
2. Decode/ Generate control signals
3. Read source register values
4. Generate ALU outputs
5. Read or Write Data Memory for Load or Store instruction
6. New PC value is computed

Update destination register and PC at next clock edge.

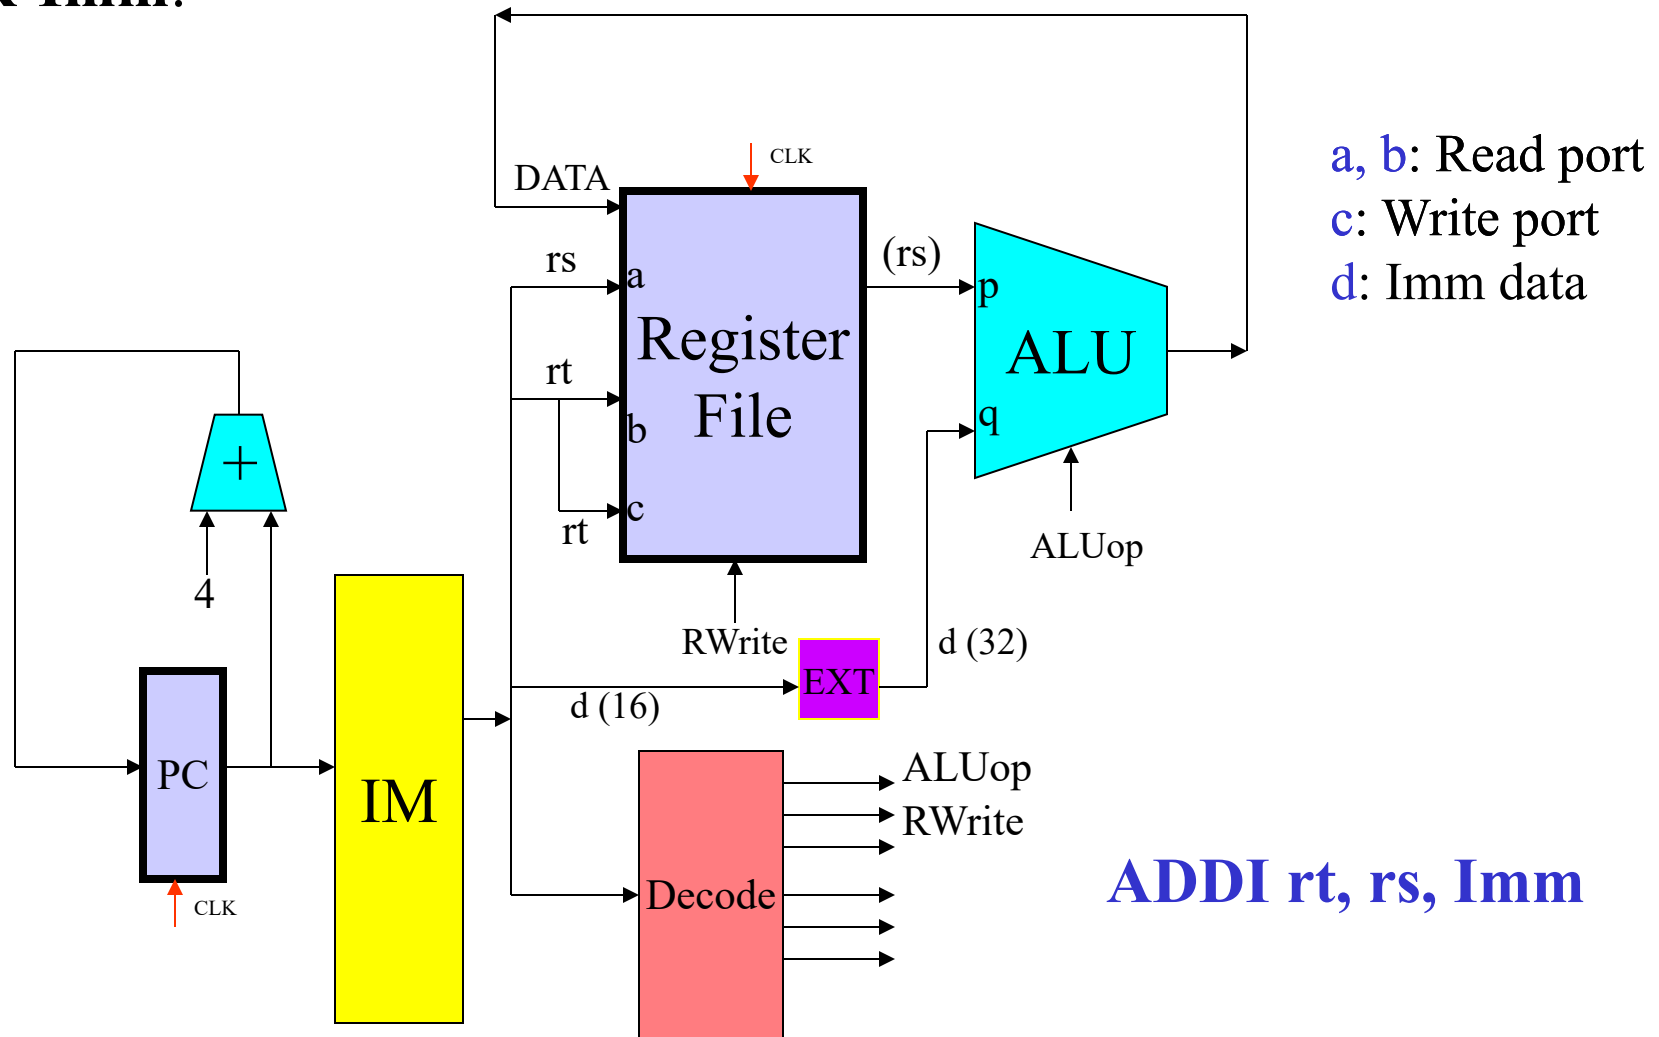
# Datapaths

**R-R Instruction:**



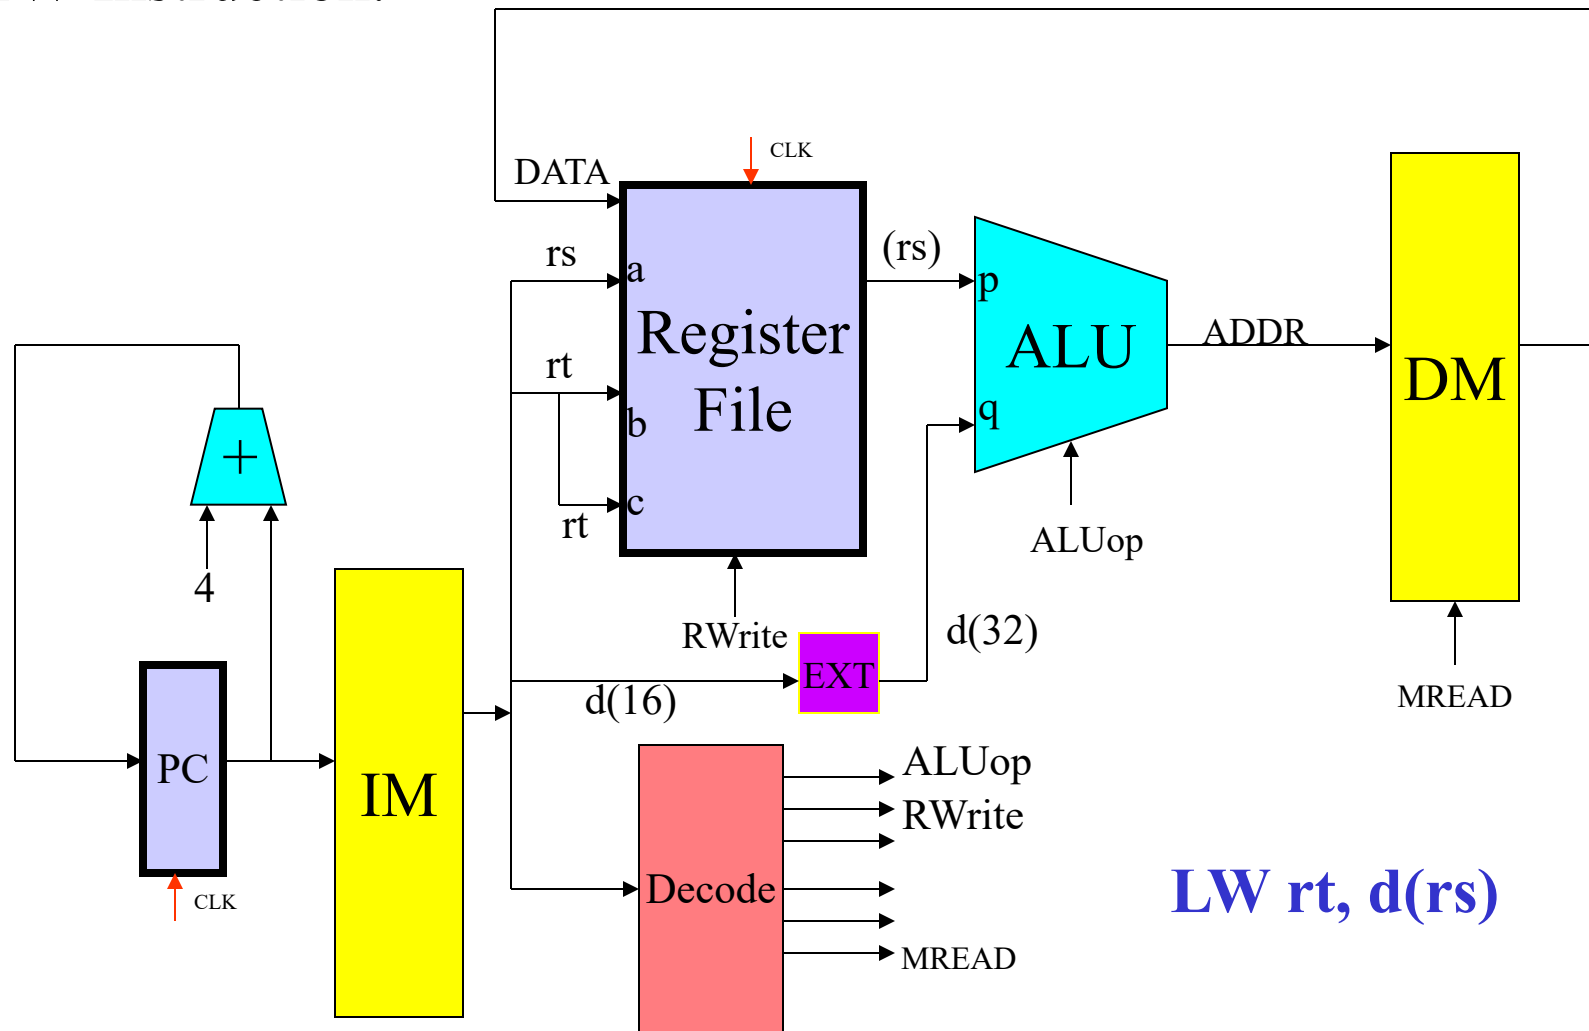
# Datapath: Register-Immediate

**R-Imm:**



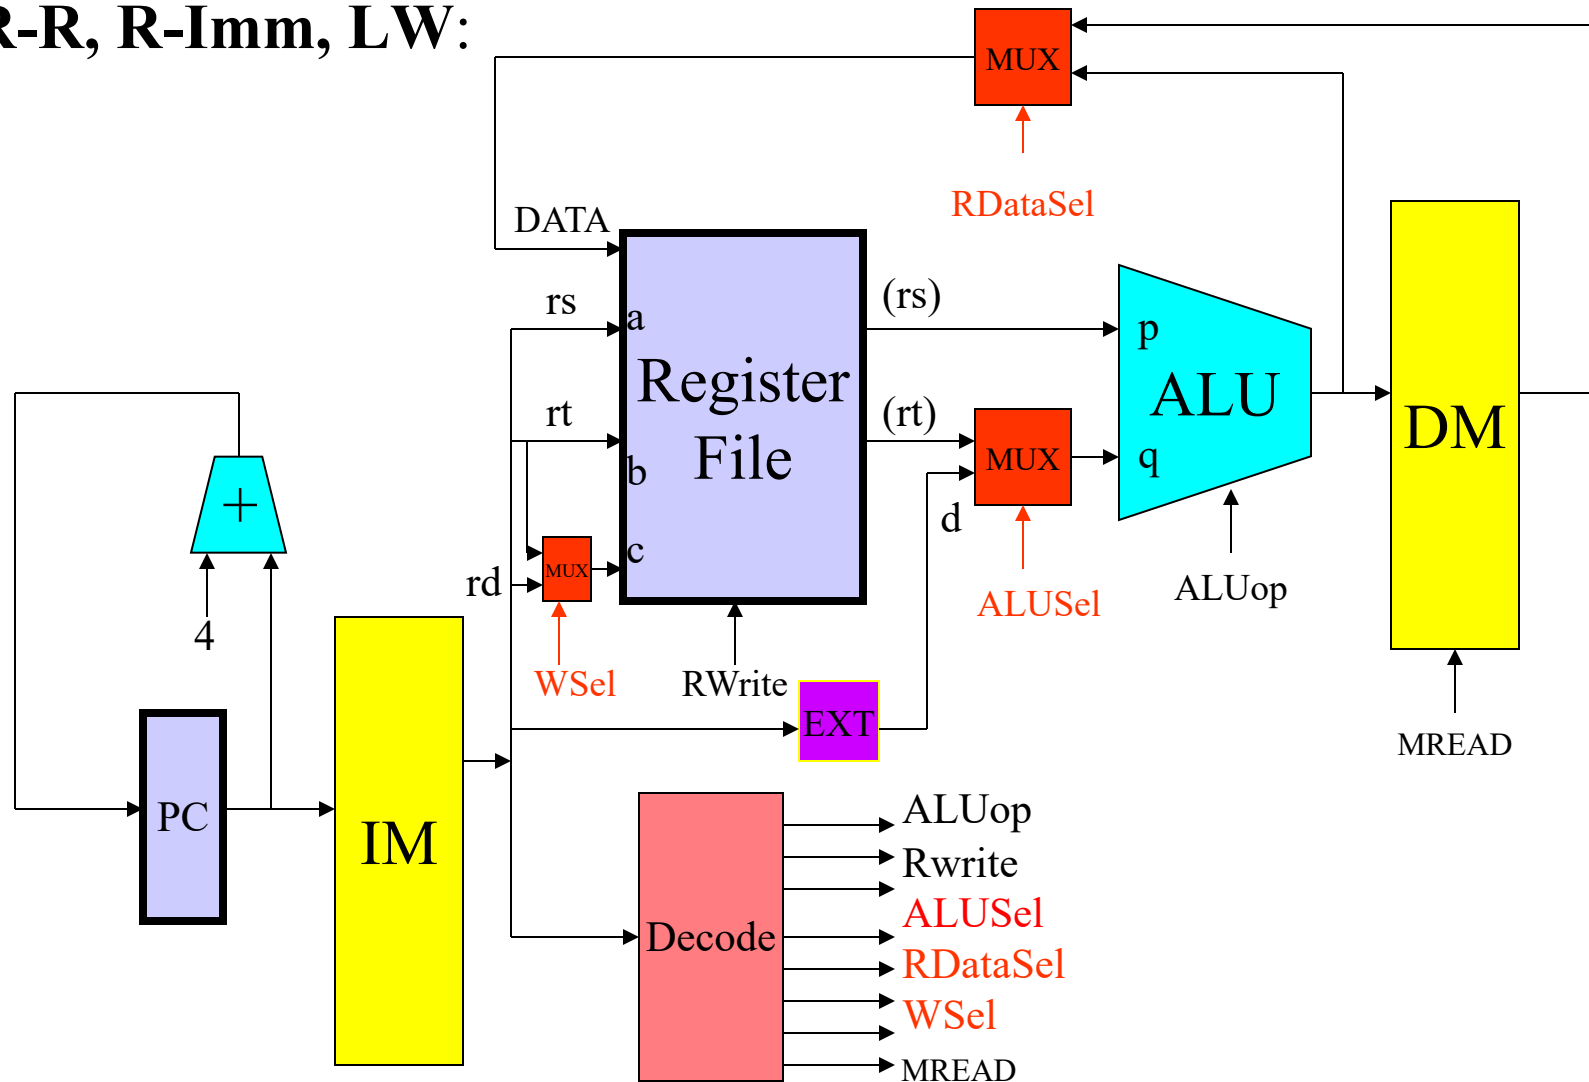
# Datapaths

**LW Instruction:**



# Datapaths

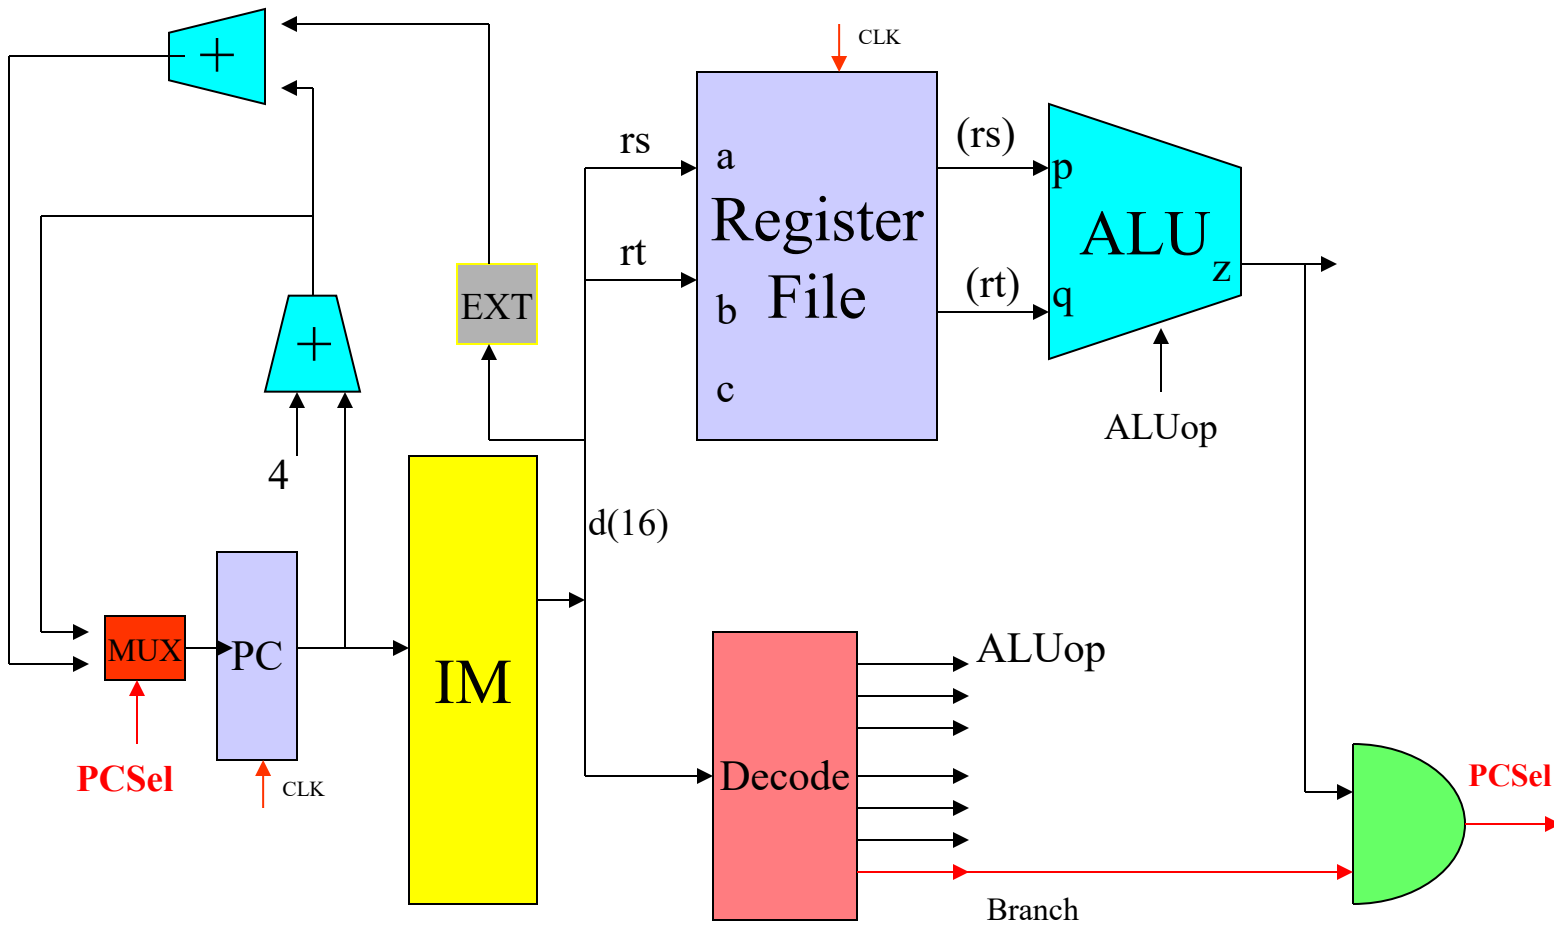
**R-R, R-Imm, LW:**



# Datapaths

**BEQ** Instruction:

**BEQ** rs, rt, d





## Performance Model

Instruction Throughput (MIPS rating):

- Clock Frequency (cycles/sec) x IPC (instructions/cycle) x  $10^{-6}$

CPI (Cycles per Instruction)

Measure by counting instruction completion rate

Single-cycle design: Each instruction takes 1 clock cycle (CPI = 1)

IPC (Instructions/Cycle) =  $1/\text{CPI}$

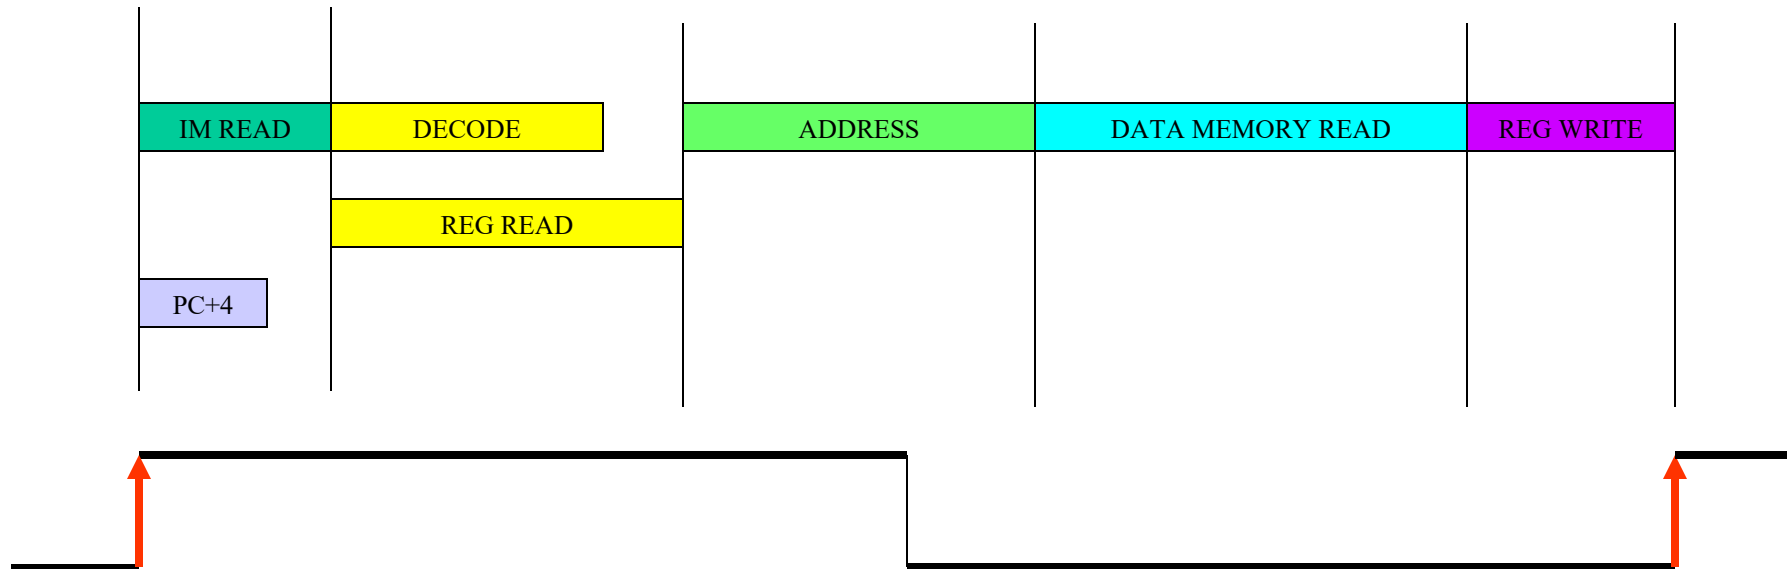
How do we improve performance?

- Increase Frequency
- Increase IPC
  - Workload to measure IPC?
  - SPEC Benchmarks


# Single Cycle Design

- Cycle time determined by longest instruction
- No reuse of Functional Modules (Separate IM and DM, ALU etc)

LW



ADD



Blank Slide

# Pipelined Processor Design

Instruction execution decomposed into 5 stages

IF: Instruction Fetch

ID/REG: Instruction Decode and Register Read

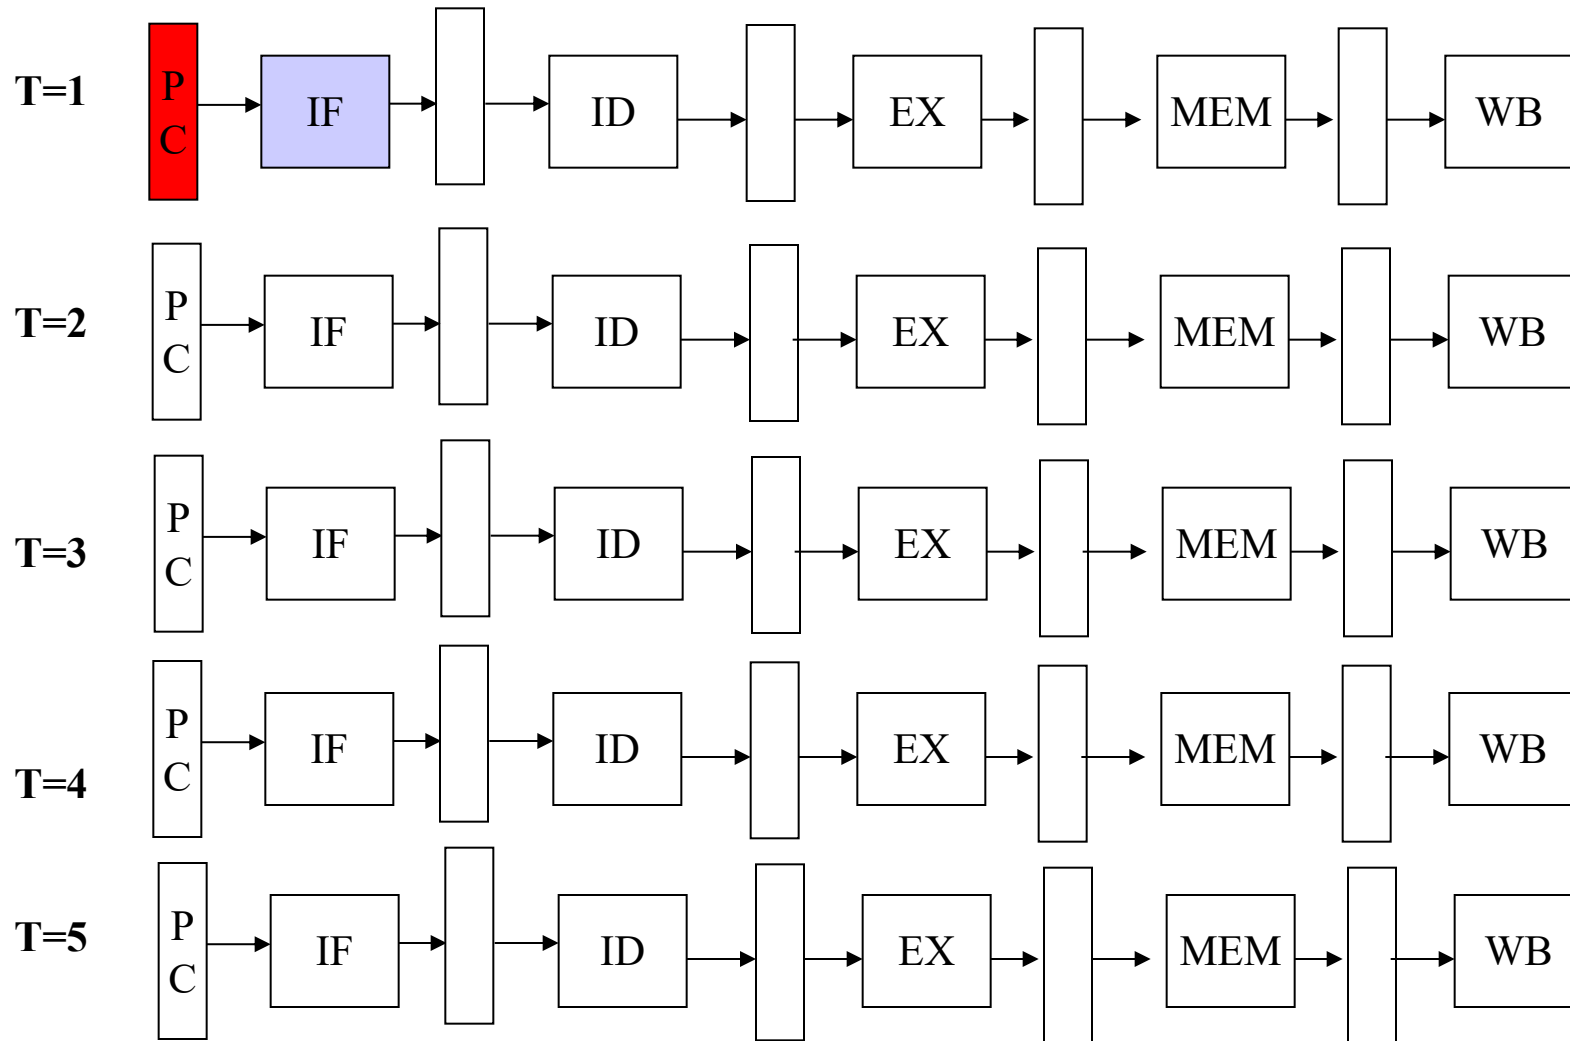
ALU: ALU operation

MEM: Data memory Read or Write

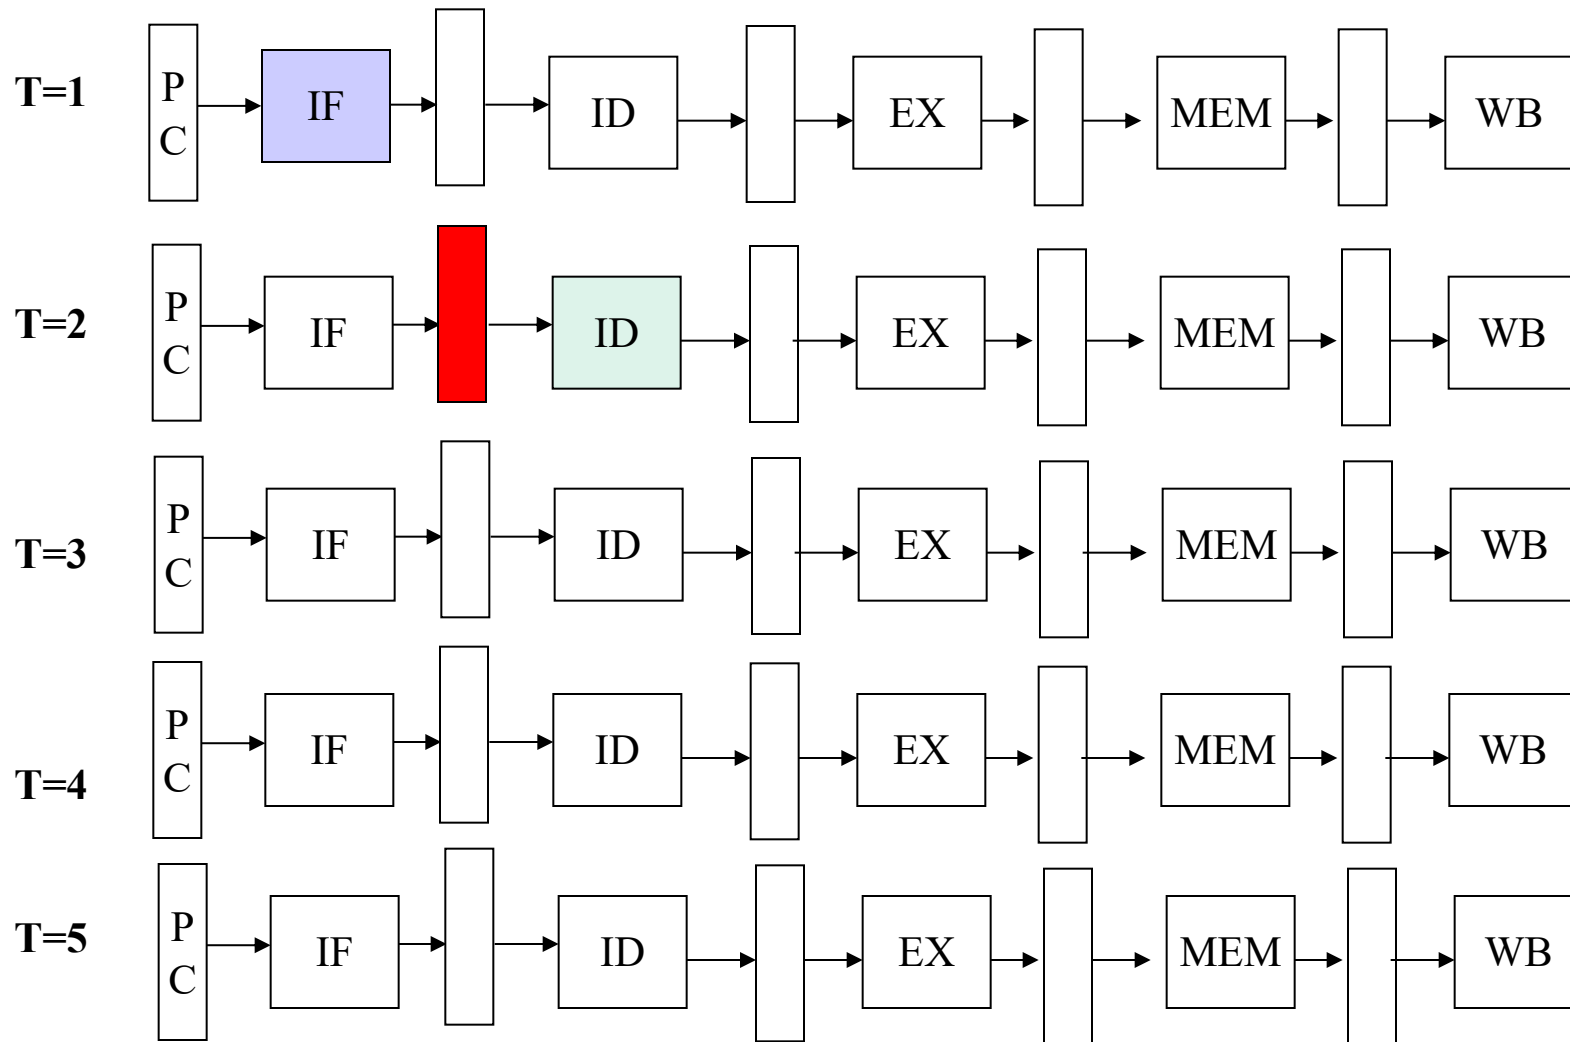
WB: Writeback to Register

- Successive stages of the pipeline separated by Pipeline Registers
- Each stage takes 1 clock cycle
  - Inputs read from the pipeline register on the left
  - Results clocked into pipeline register on the right

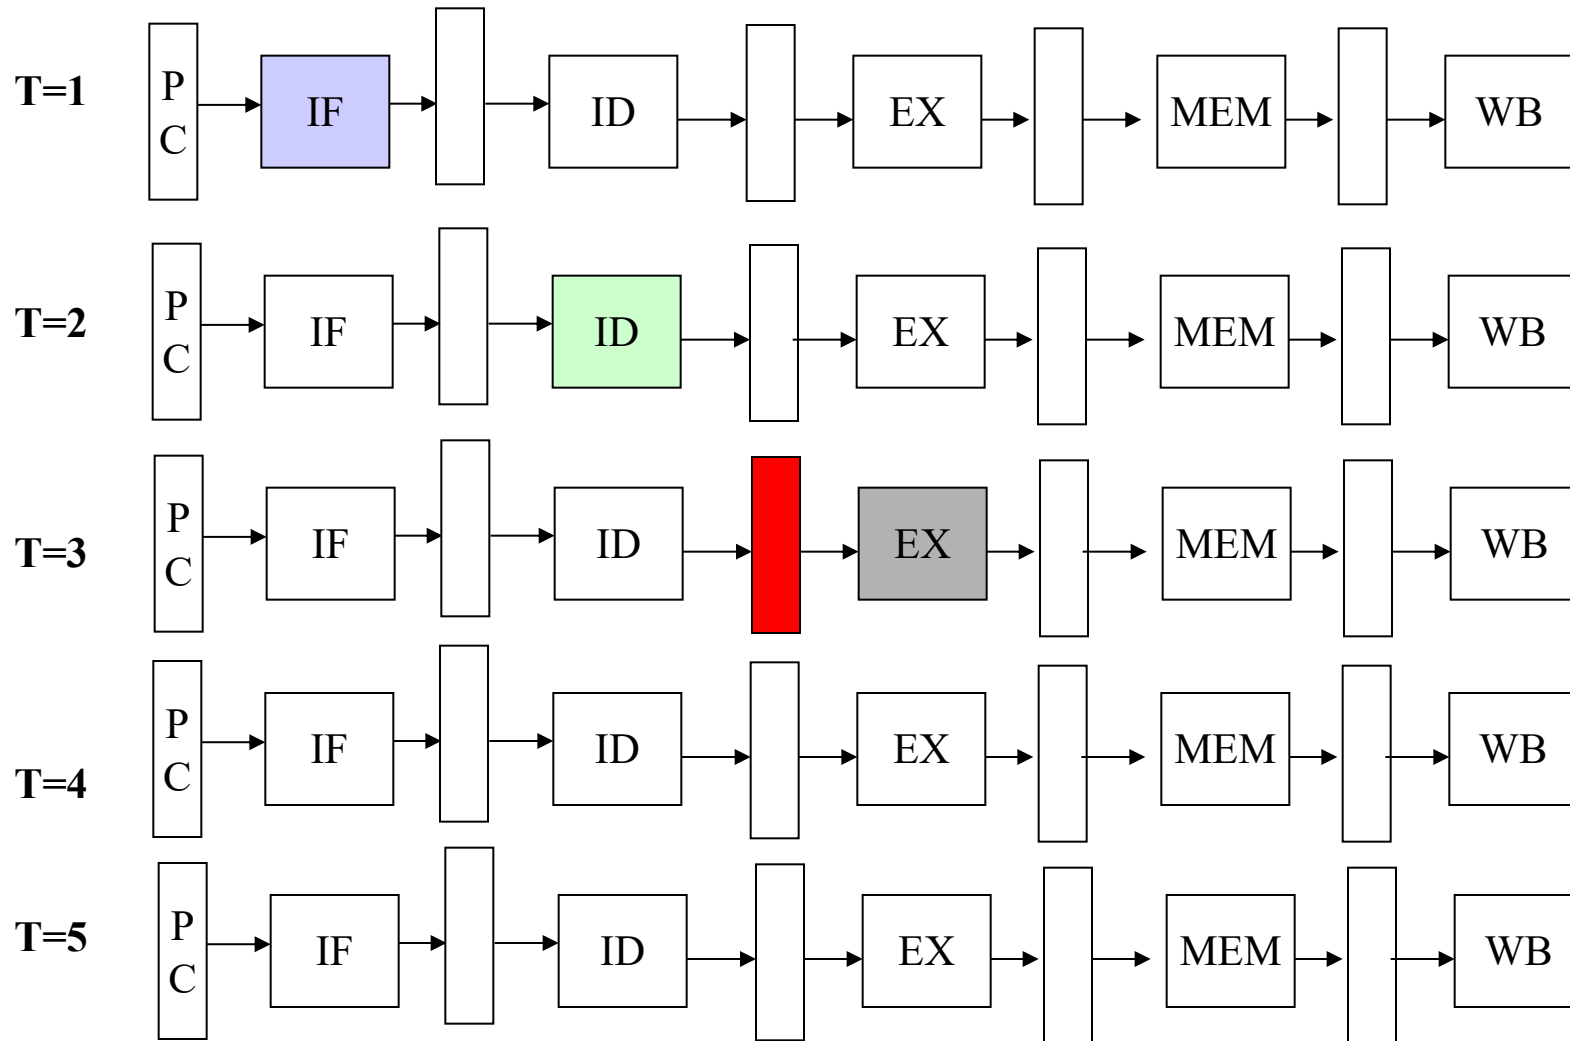
## 5-Stage Processor Pipeline



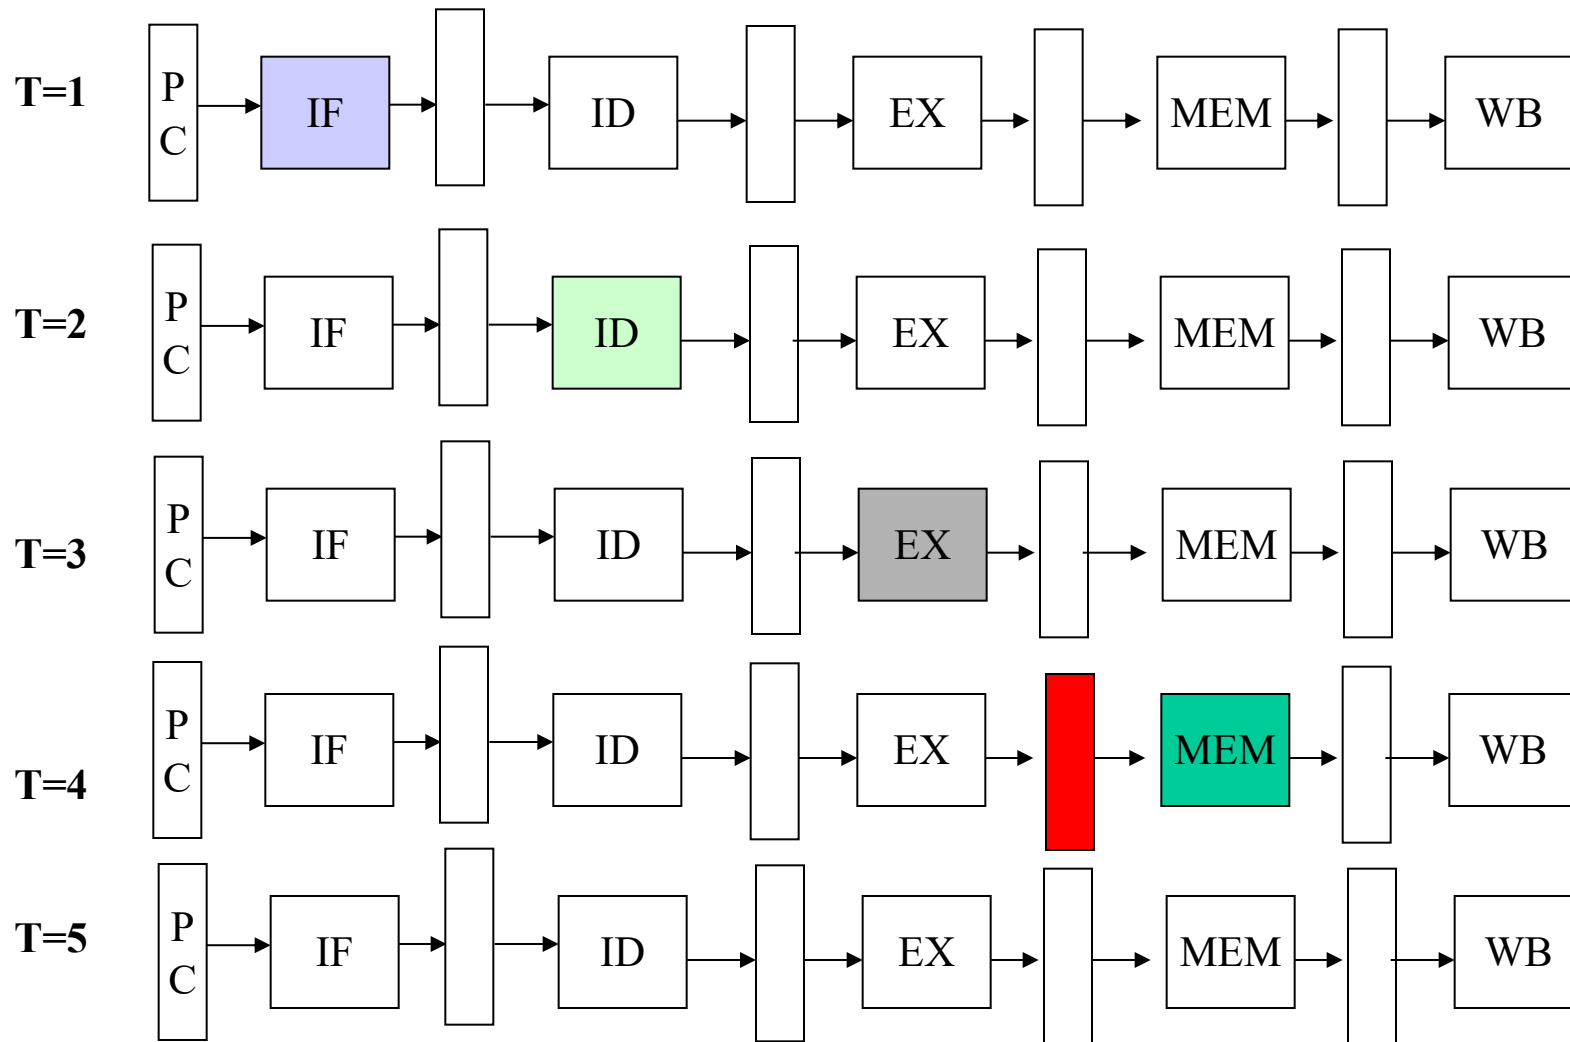
## 5-Stage Processor Pipeline



## 5-Stage Processor Pipeline

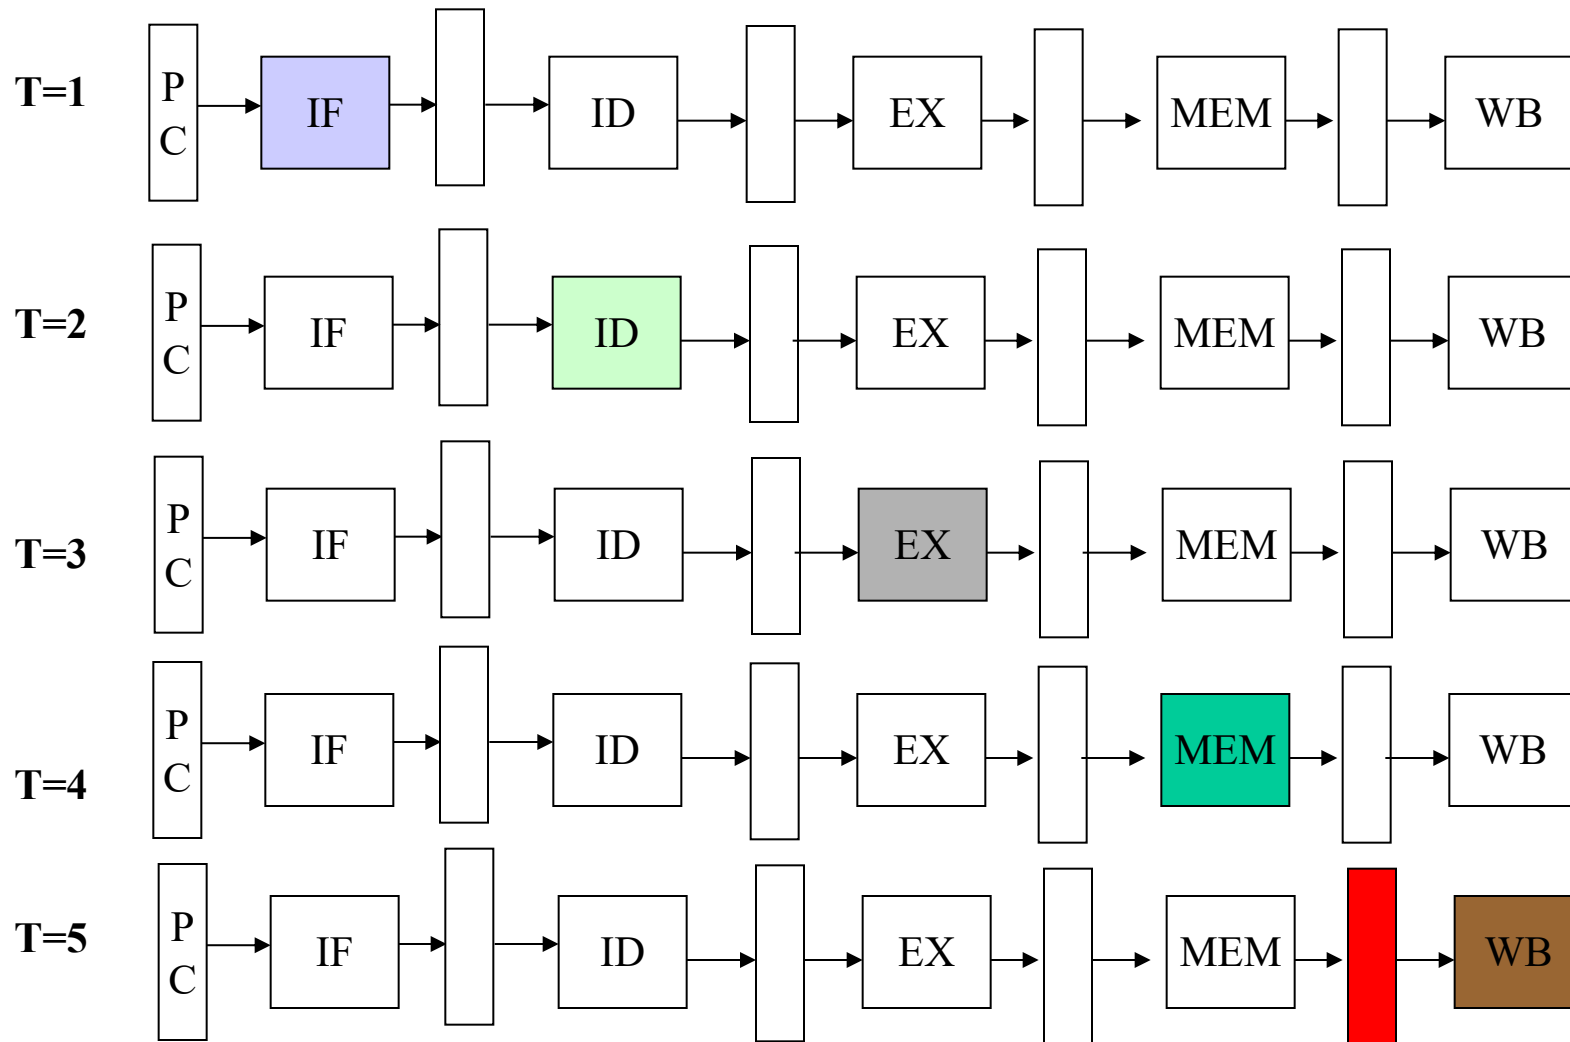


## 5-Stage Processor Pipeline






## 5-Stage Processor Pipeline



## Execution of **LW** Instruction: Cycle 1



Pipeline Register

PC had address of current instruction

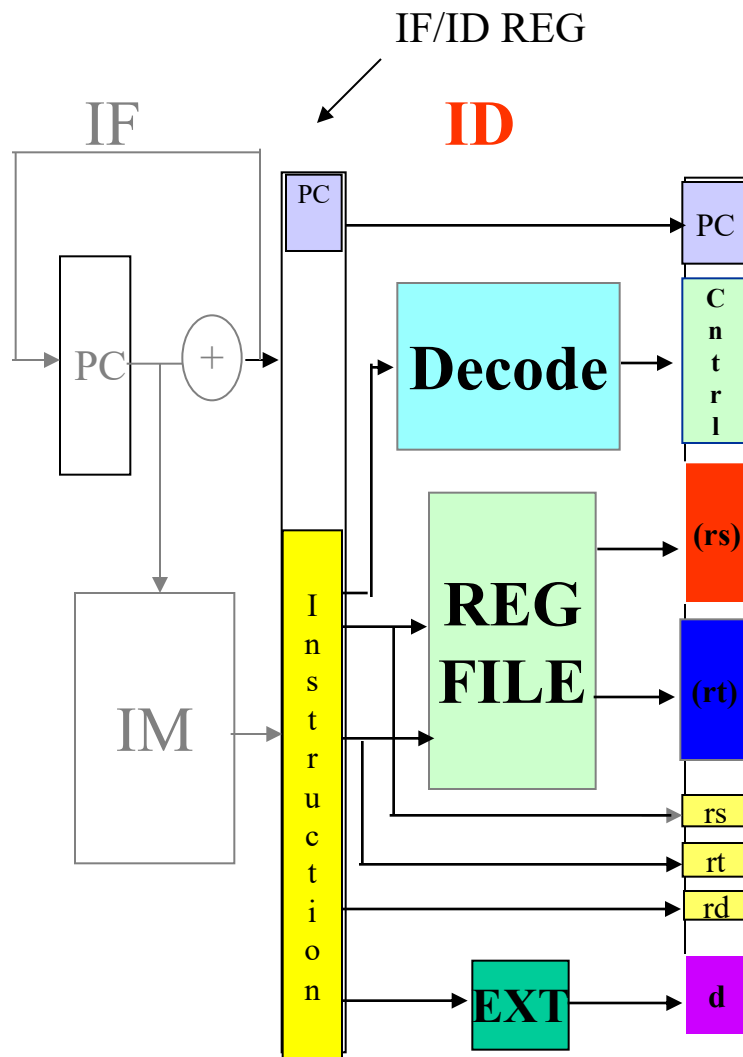
- Read instruction from Instruction Memory (IM)
- Increment PC

At clock edge:

**Write** updated PC value into **PC register** and into **Pipeline Register**

**Write** the instruction into the **Pipeline Register**

## Execution of **LW** Instruction: Cycle 2



ID/EX REG

IF/ID REG had current instruction

- Decode instruction
- Read registers

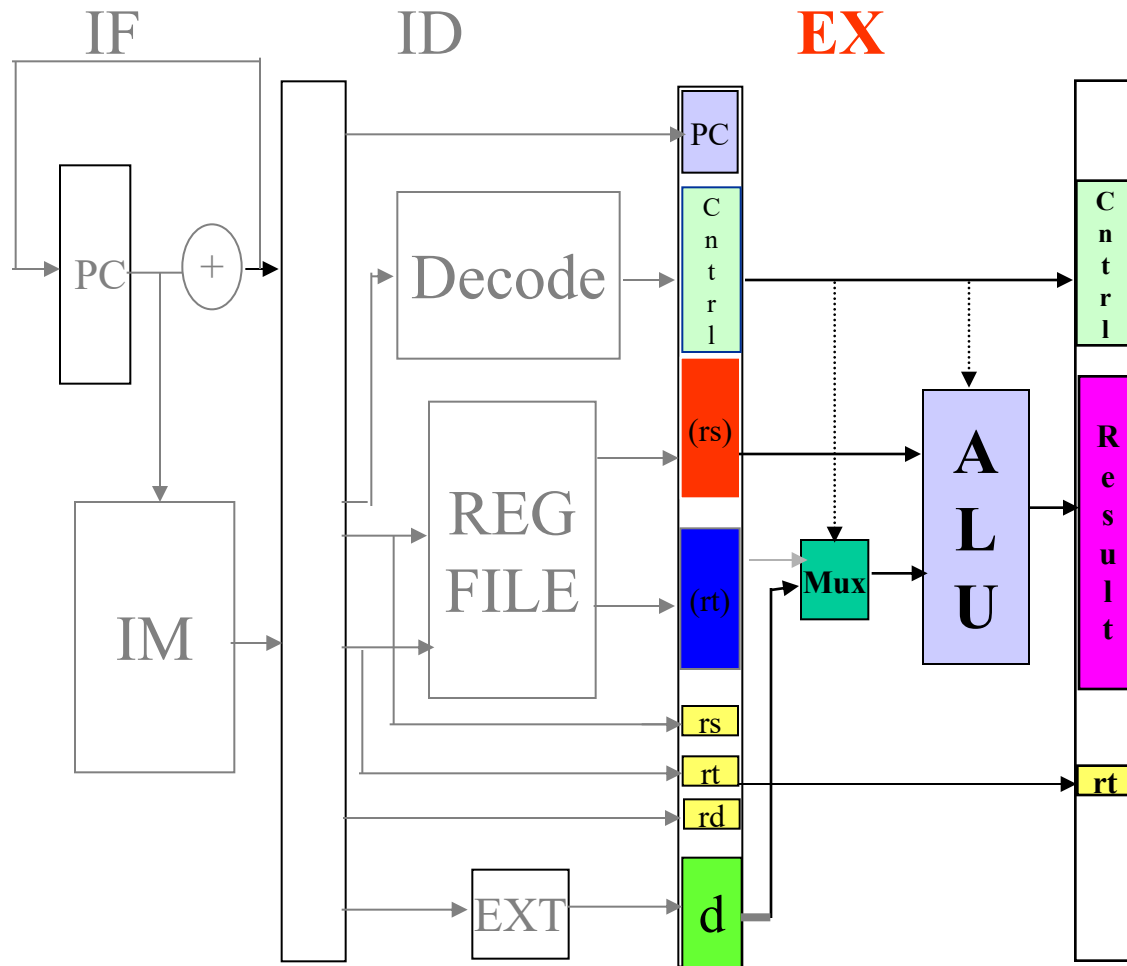
At clock edge:

**Write** decoded control signals and selected fields of instruction into **ID/EX Pipeline Register**

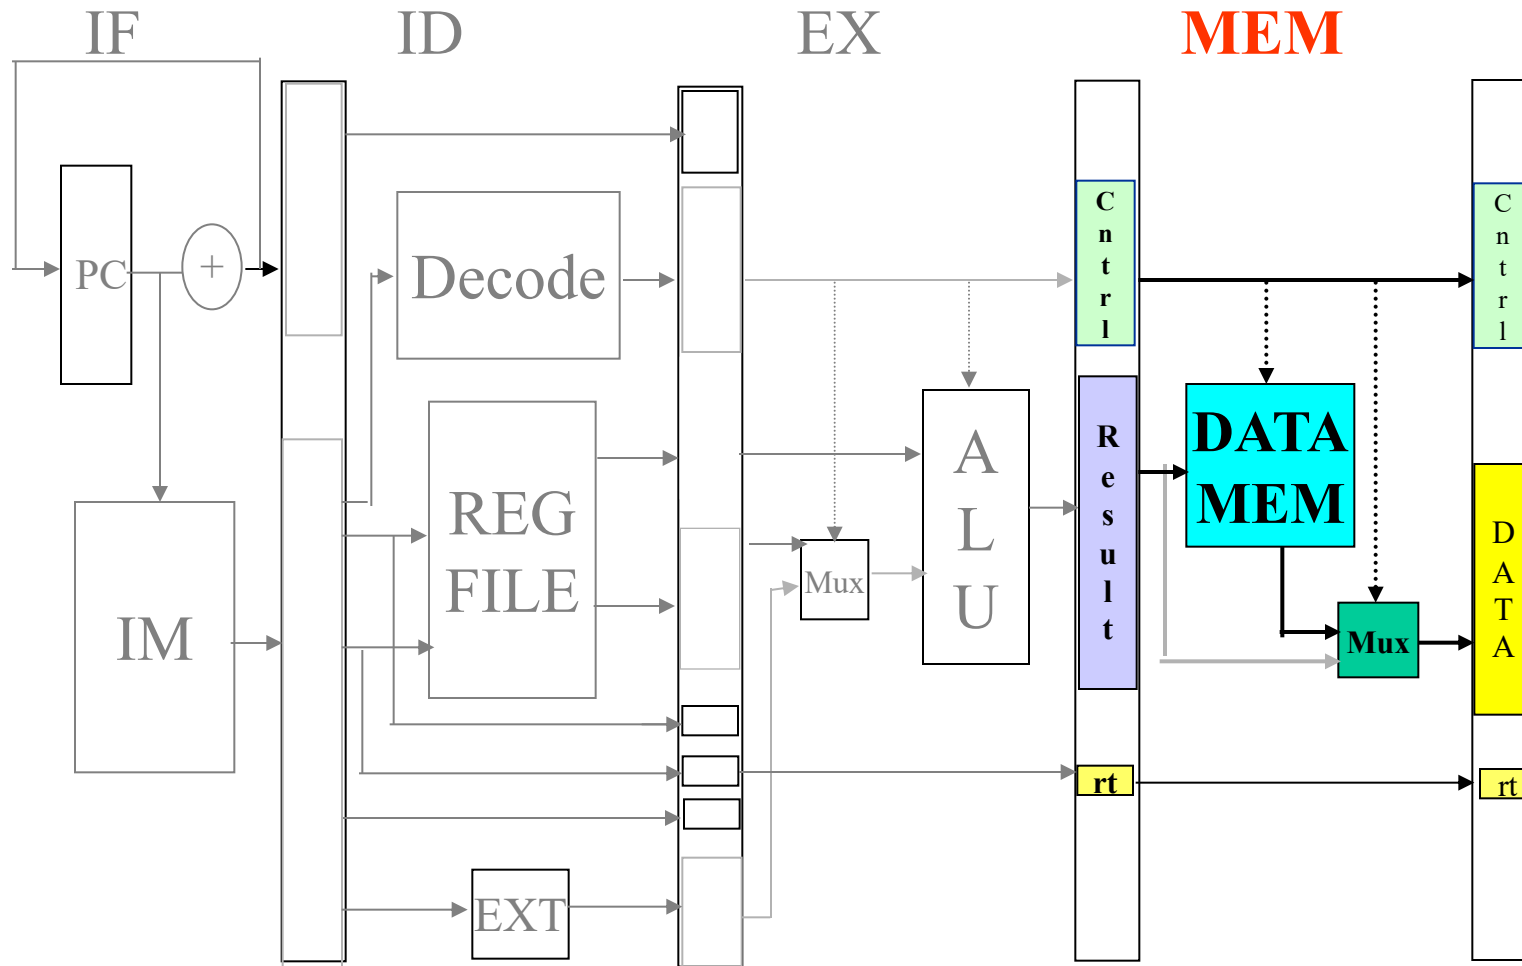
**Write** register values **rs** and **rt** into **ID/EX Pipeline Register**

**Transfer PC** value from **ID/EX** register to **ID/EX register**

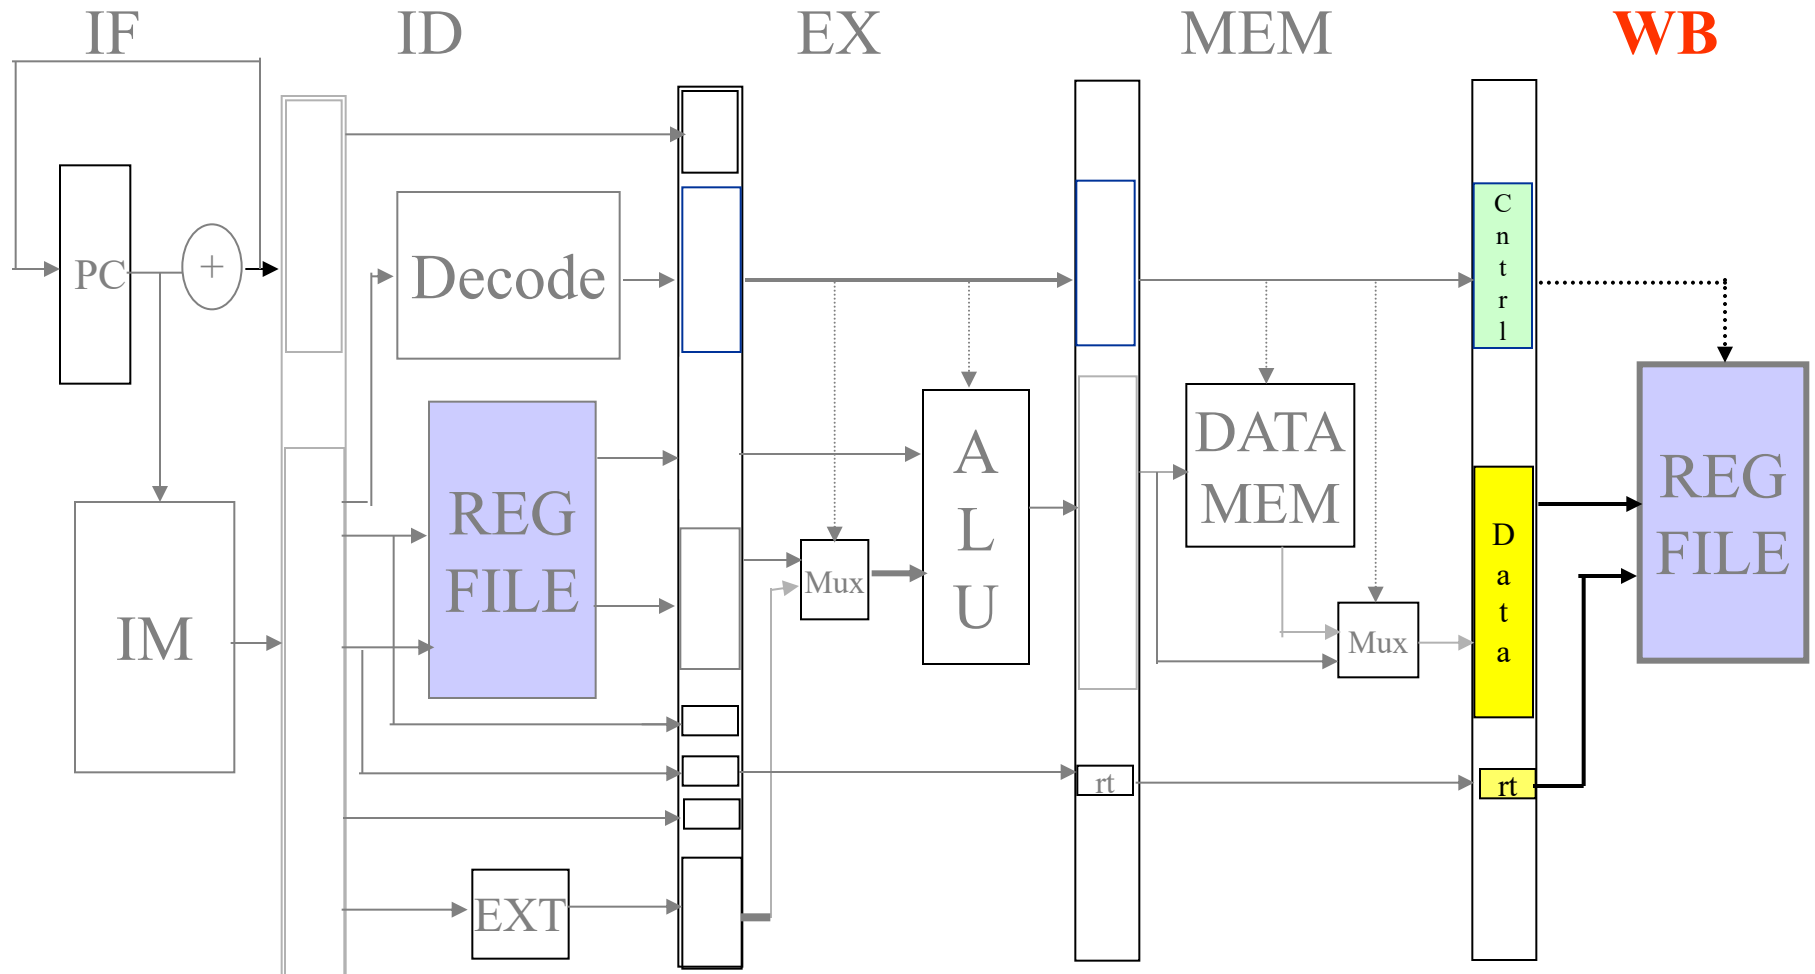
## Execution of **LW** Instruction: Cycle 3



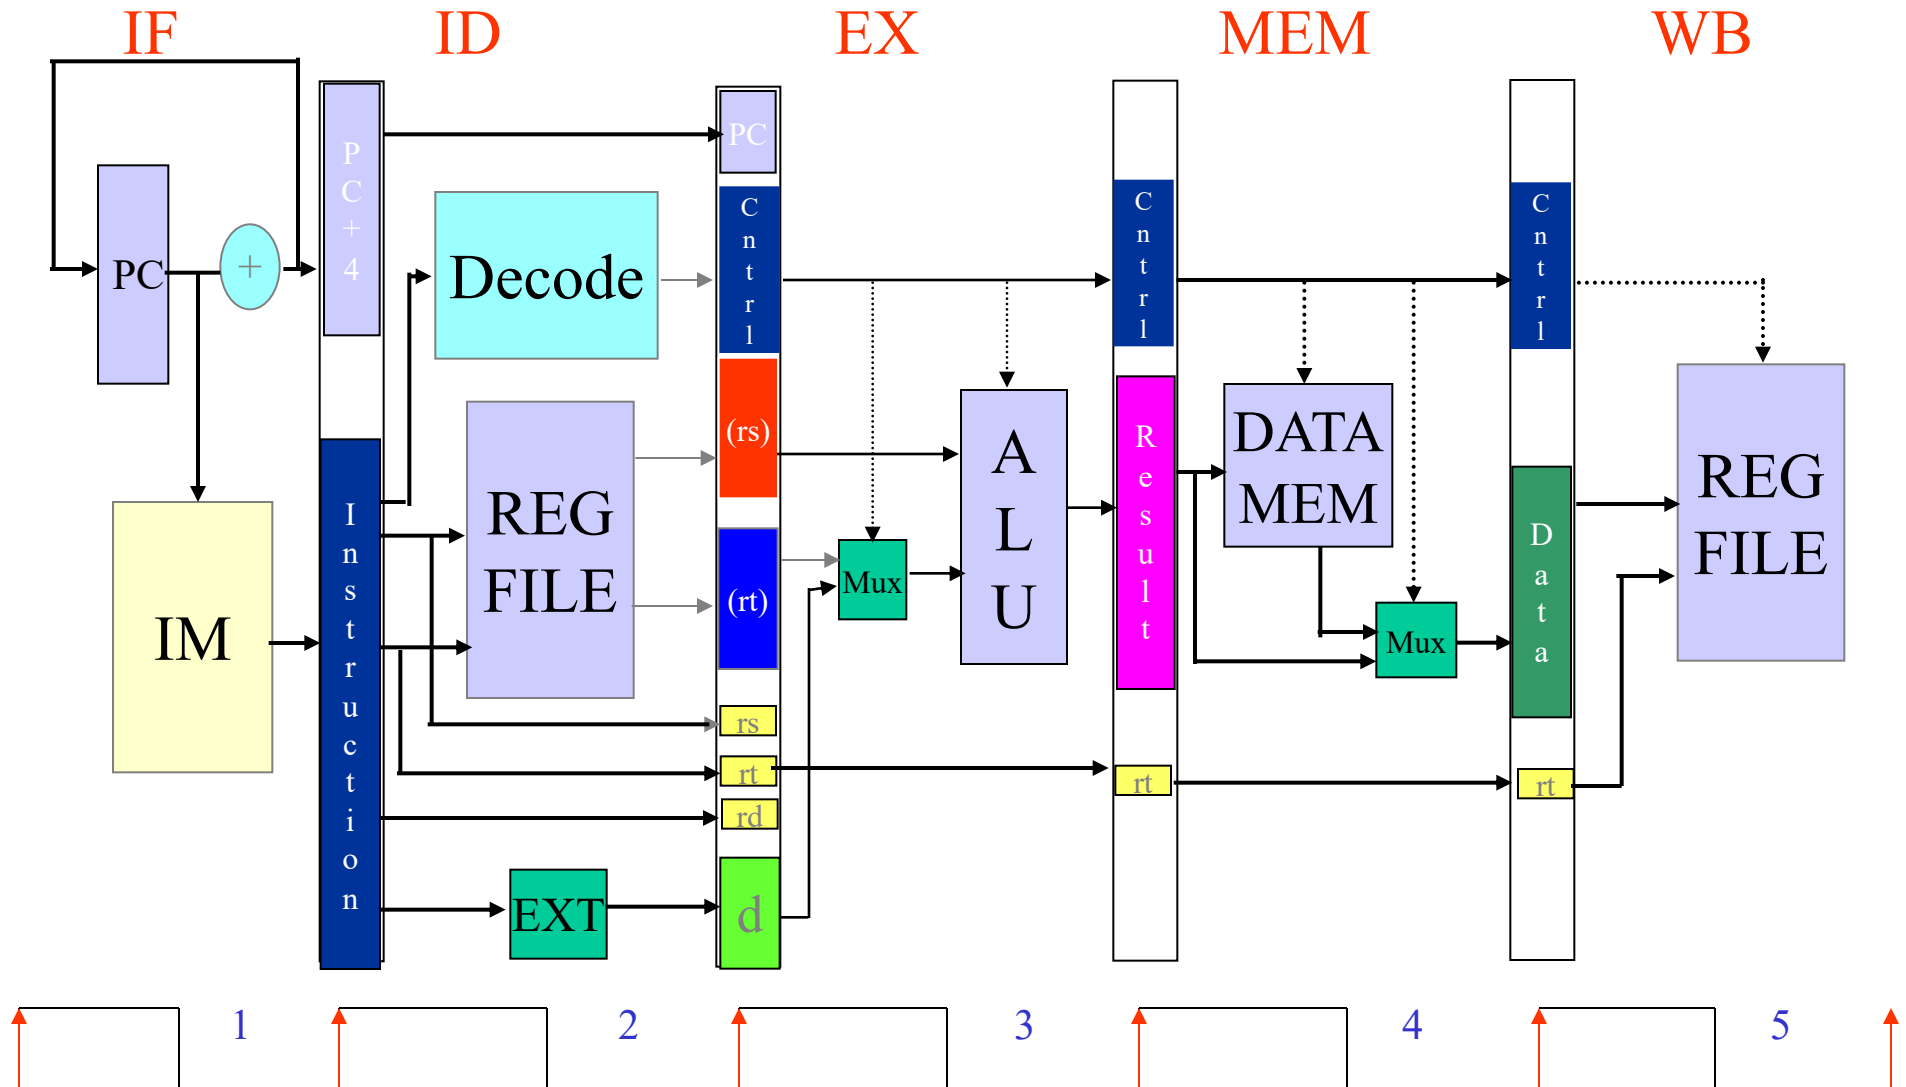
## Execution of **LW** Instruction: Cycle 4



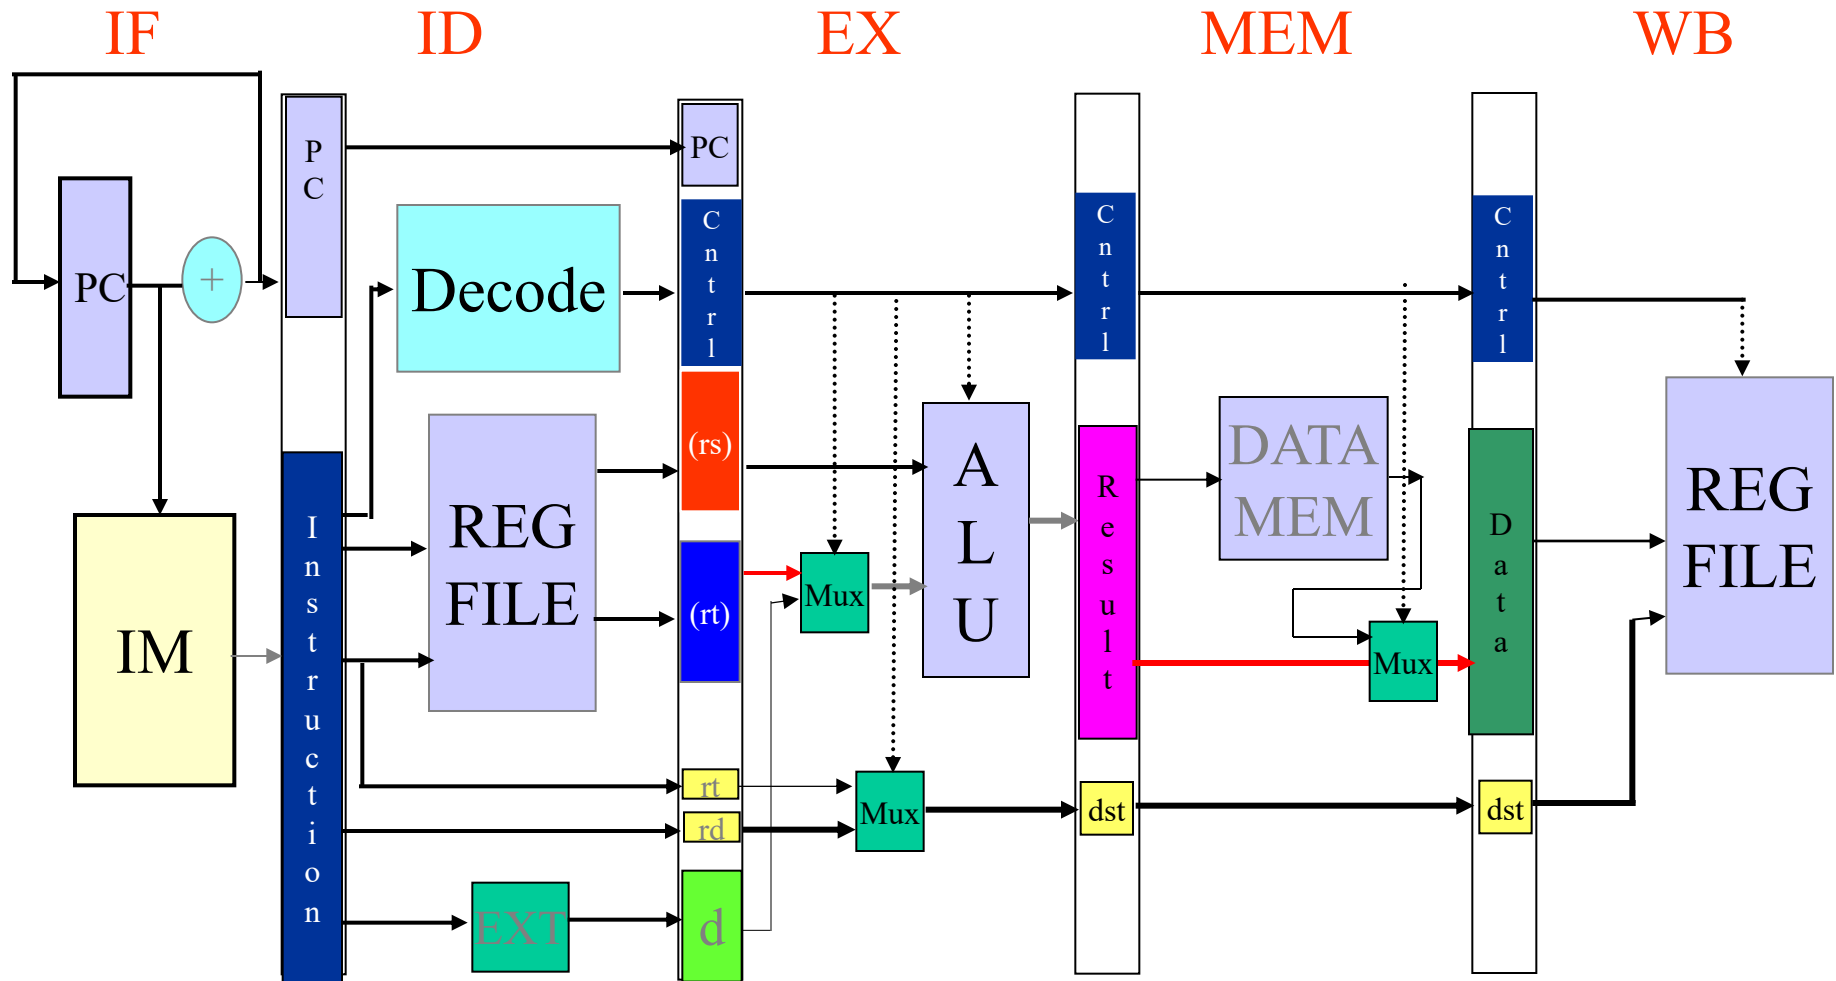
## Execution of **LW** Instruction: Cycle 5



## Execution of Load Instruction: LW rt, d(rs)



## Execution of R-R Instruction: **ADD rd, rt, rs**





# Lazy Execution of Branch Equal Instruction: BEQ rt, rs, d

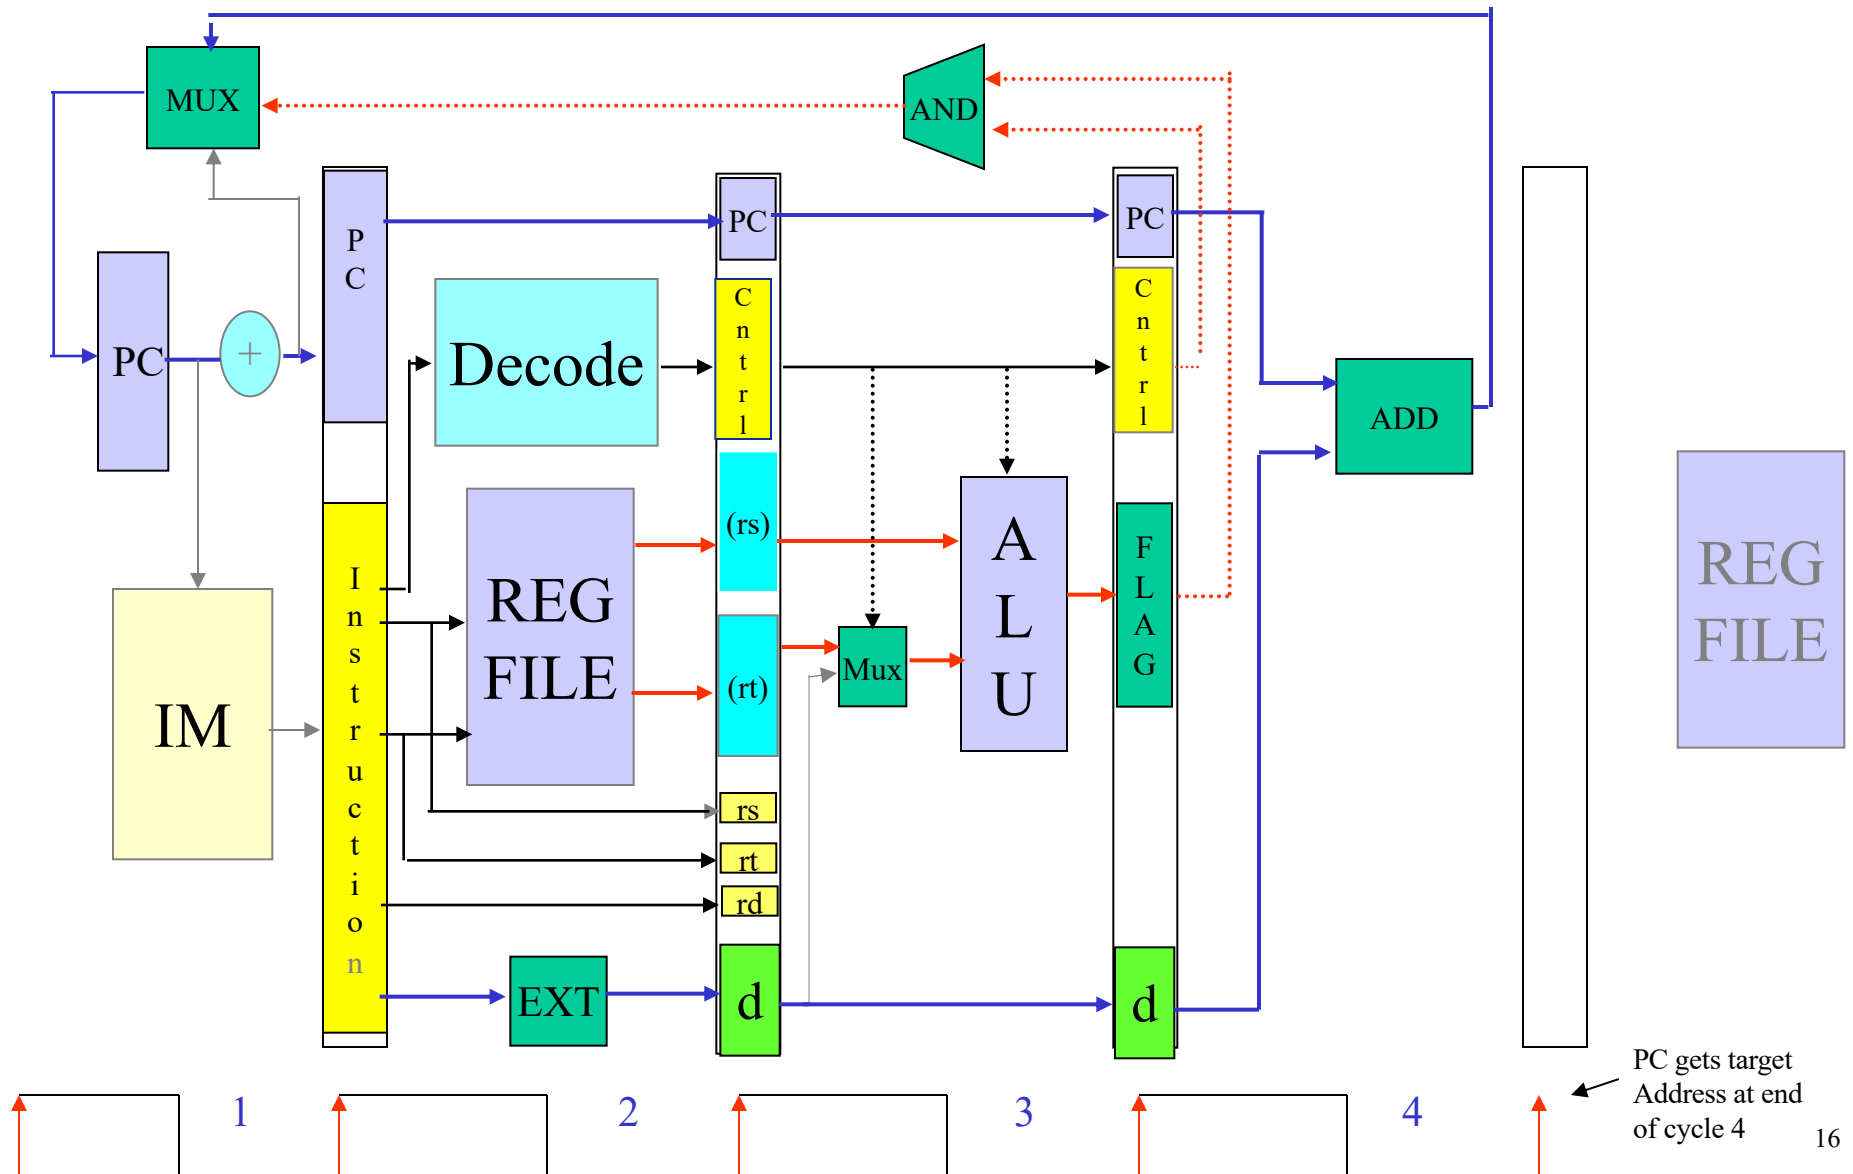
IF

ID


EX

MEM

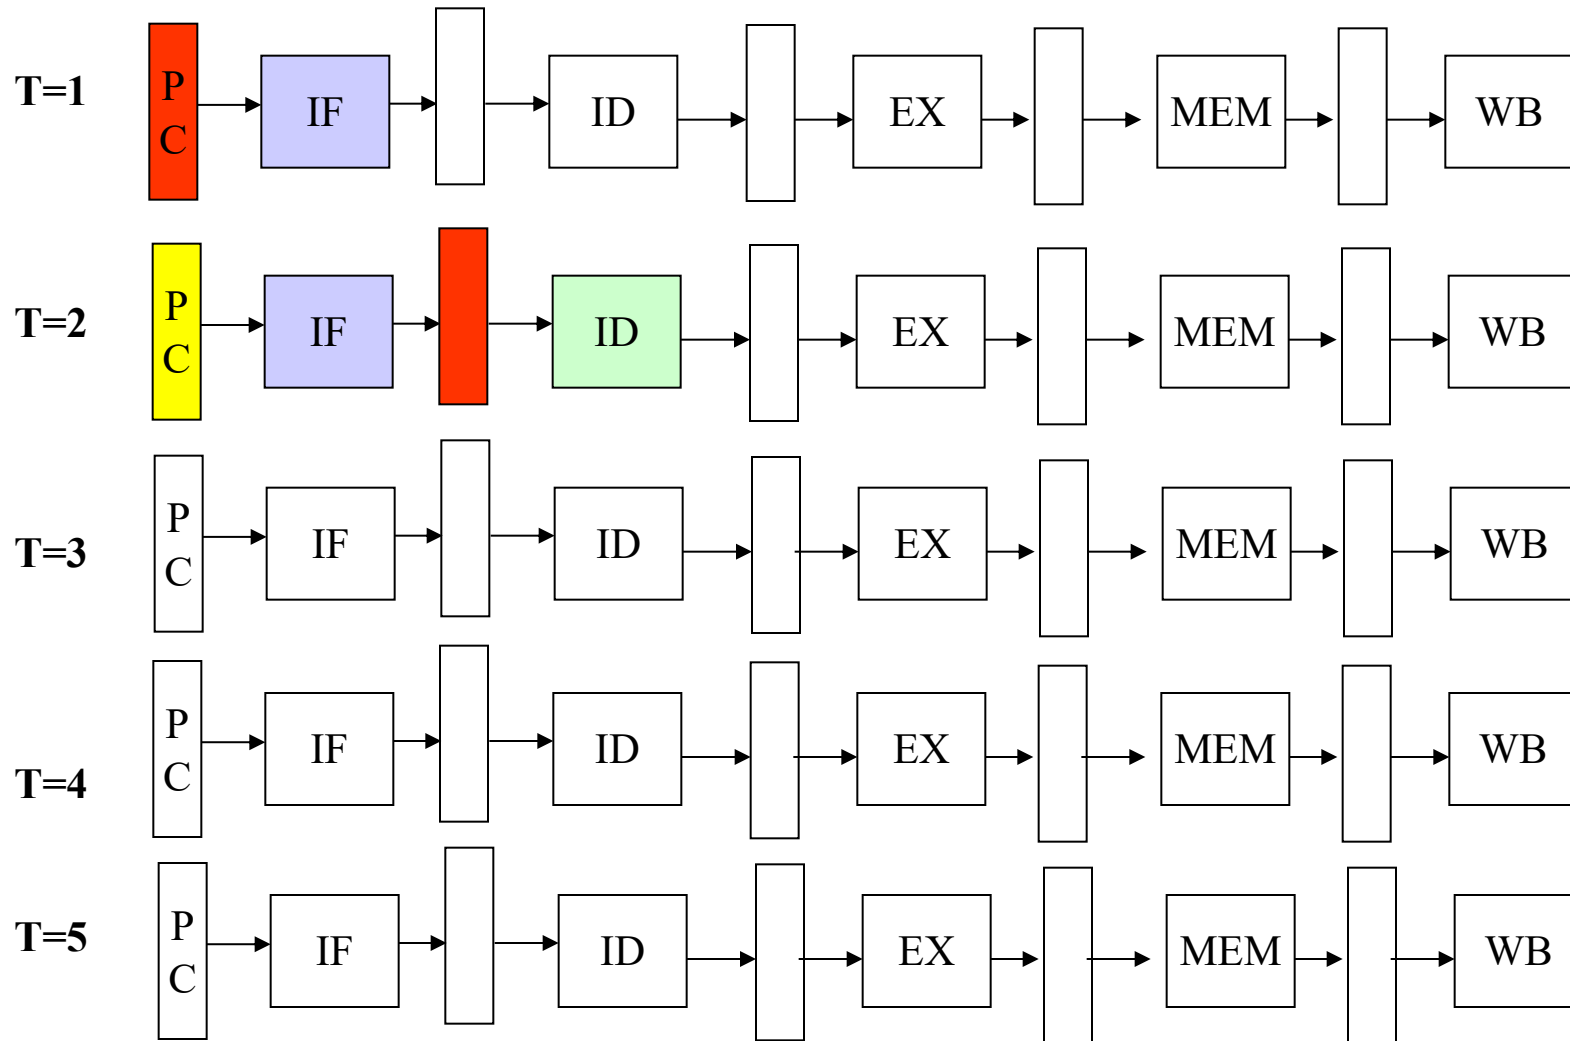
WB



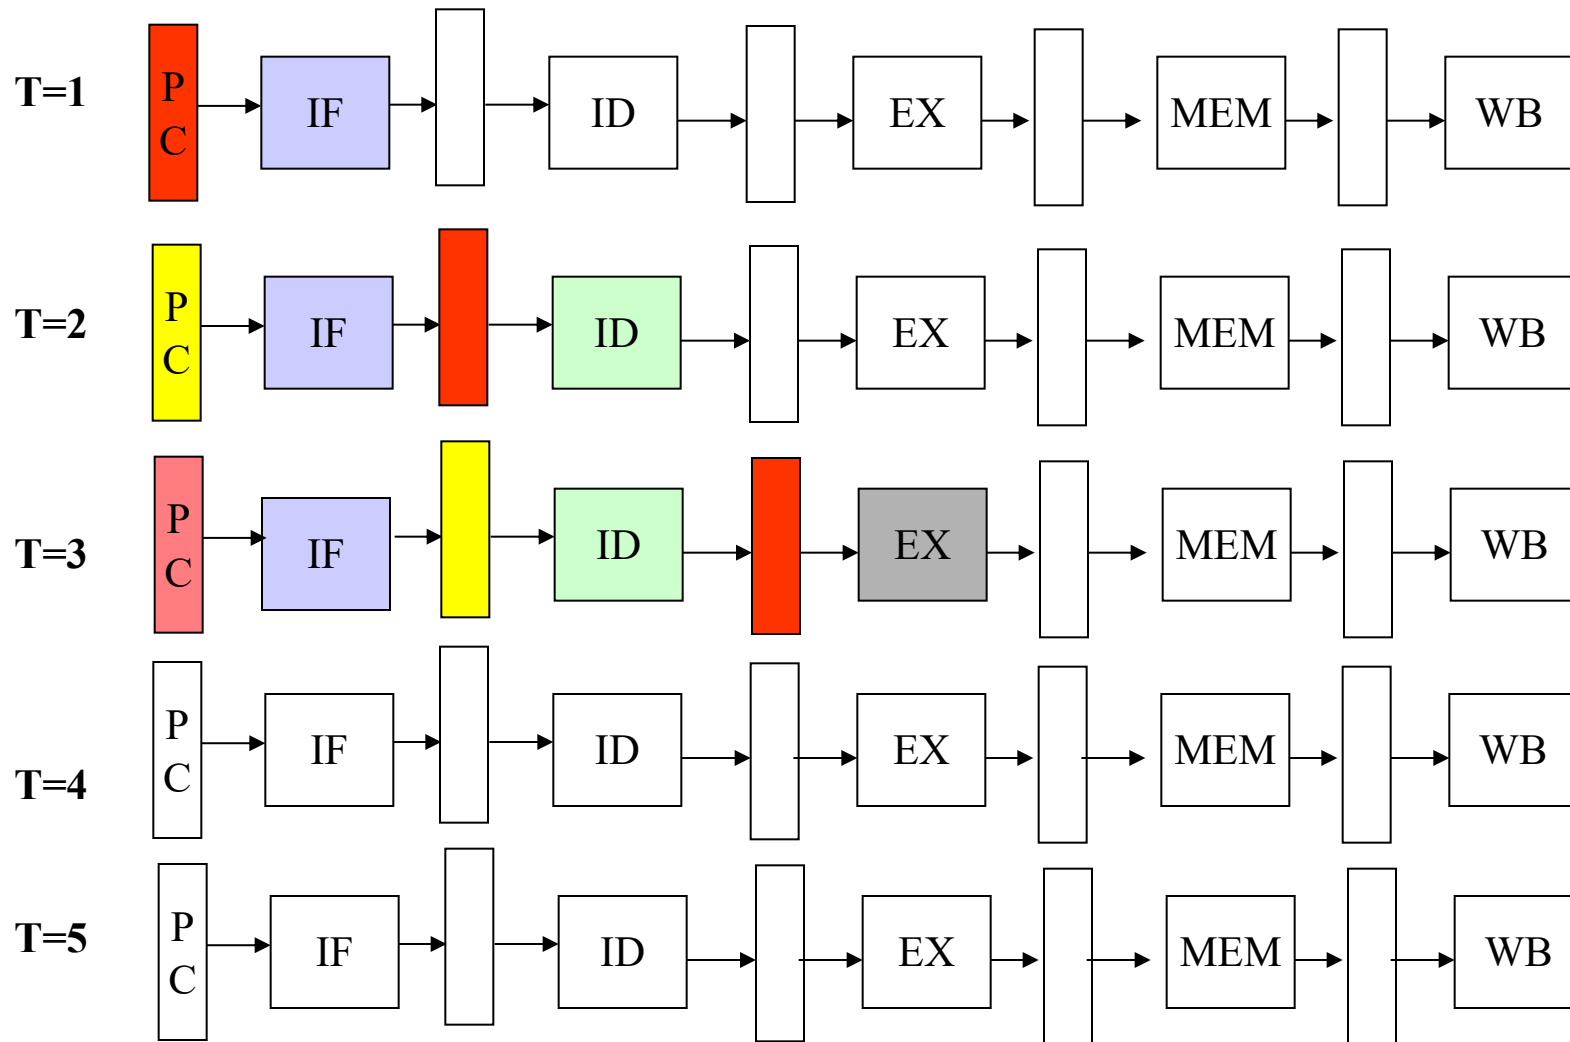
## 5-Stage Processor Pipeline



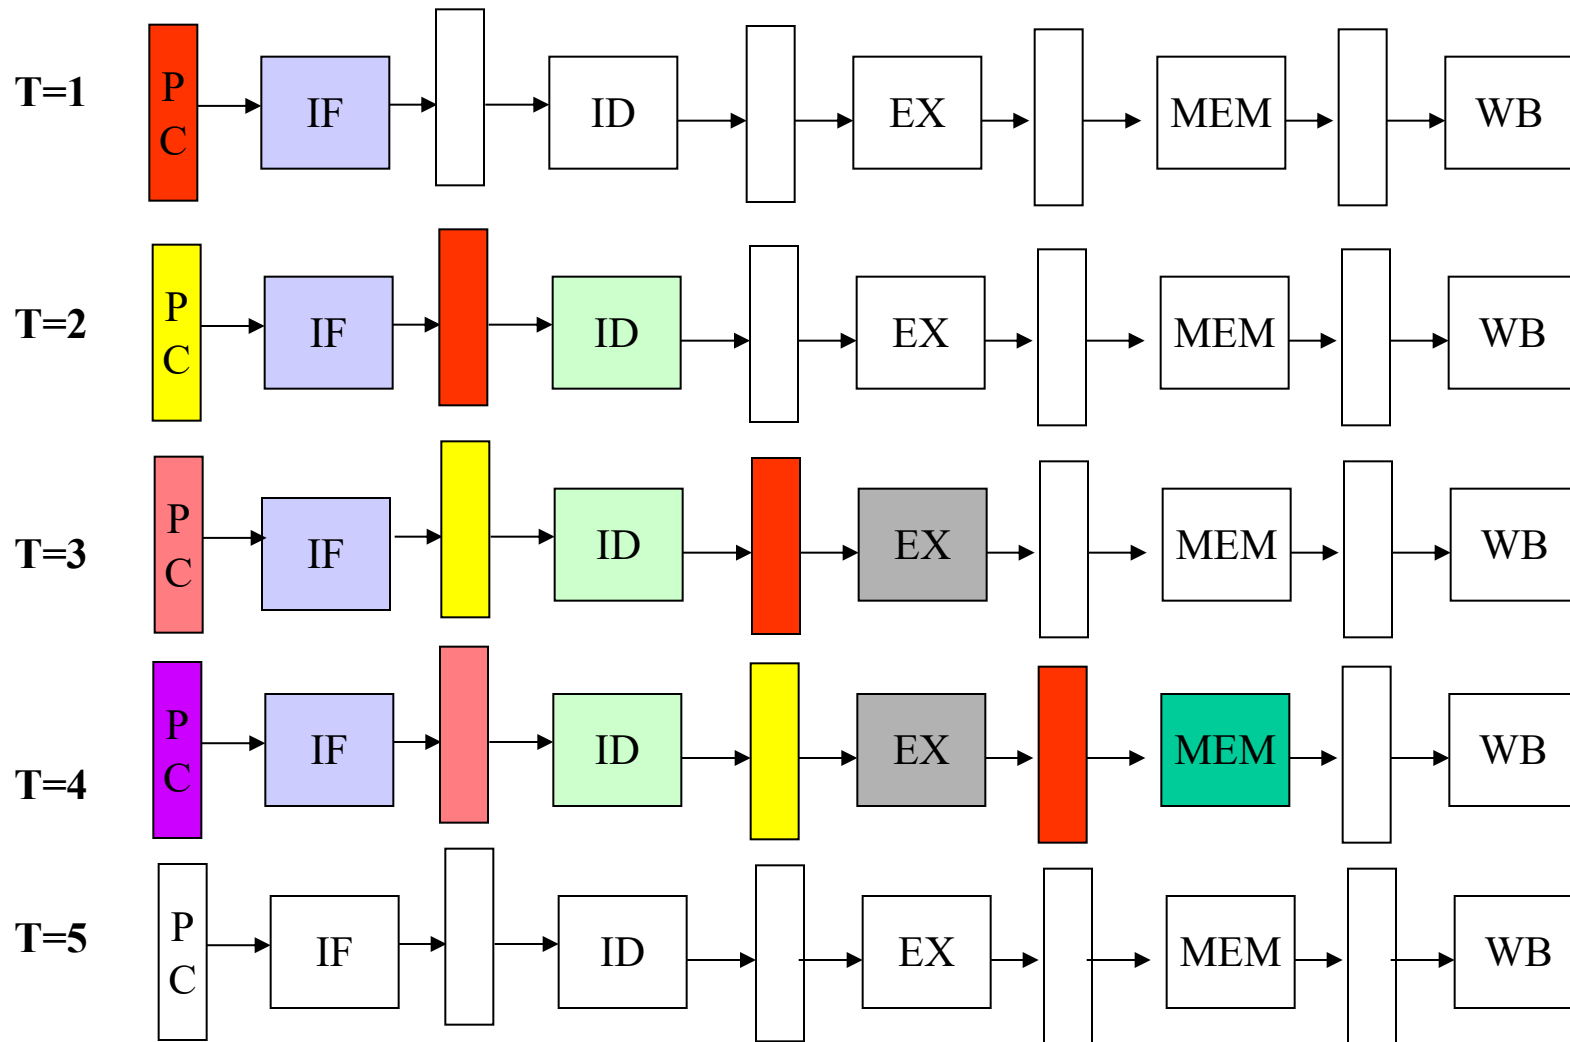
## 5-Stage Processor Pipeline



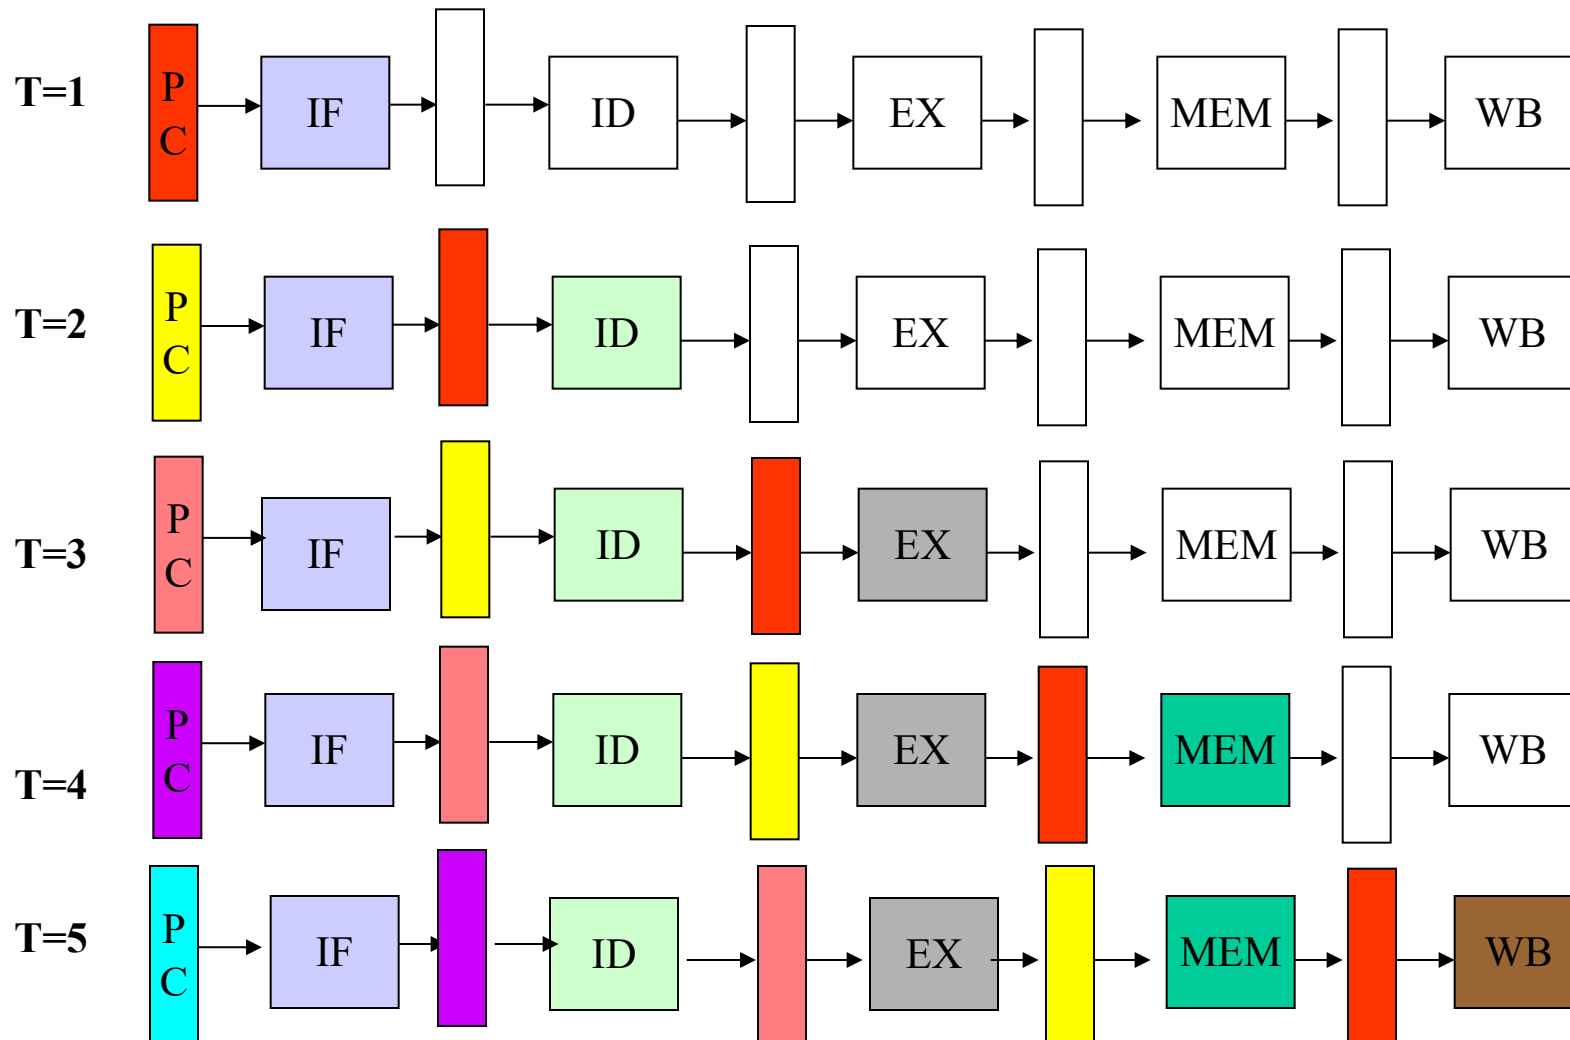
## 5-Stage Processor Pipeline



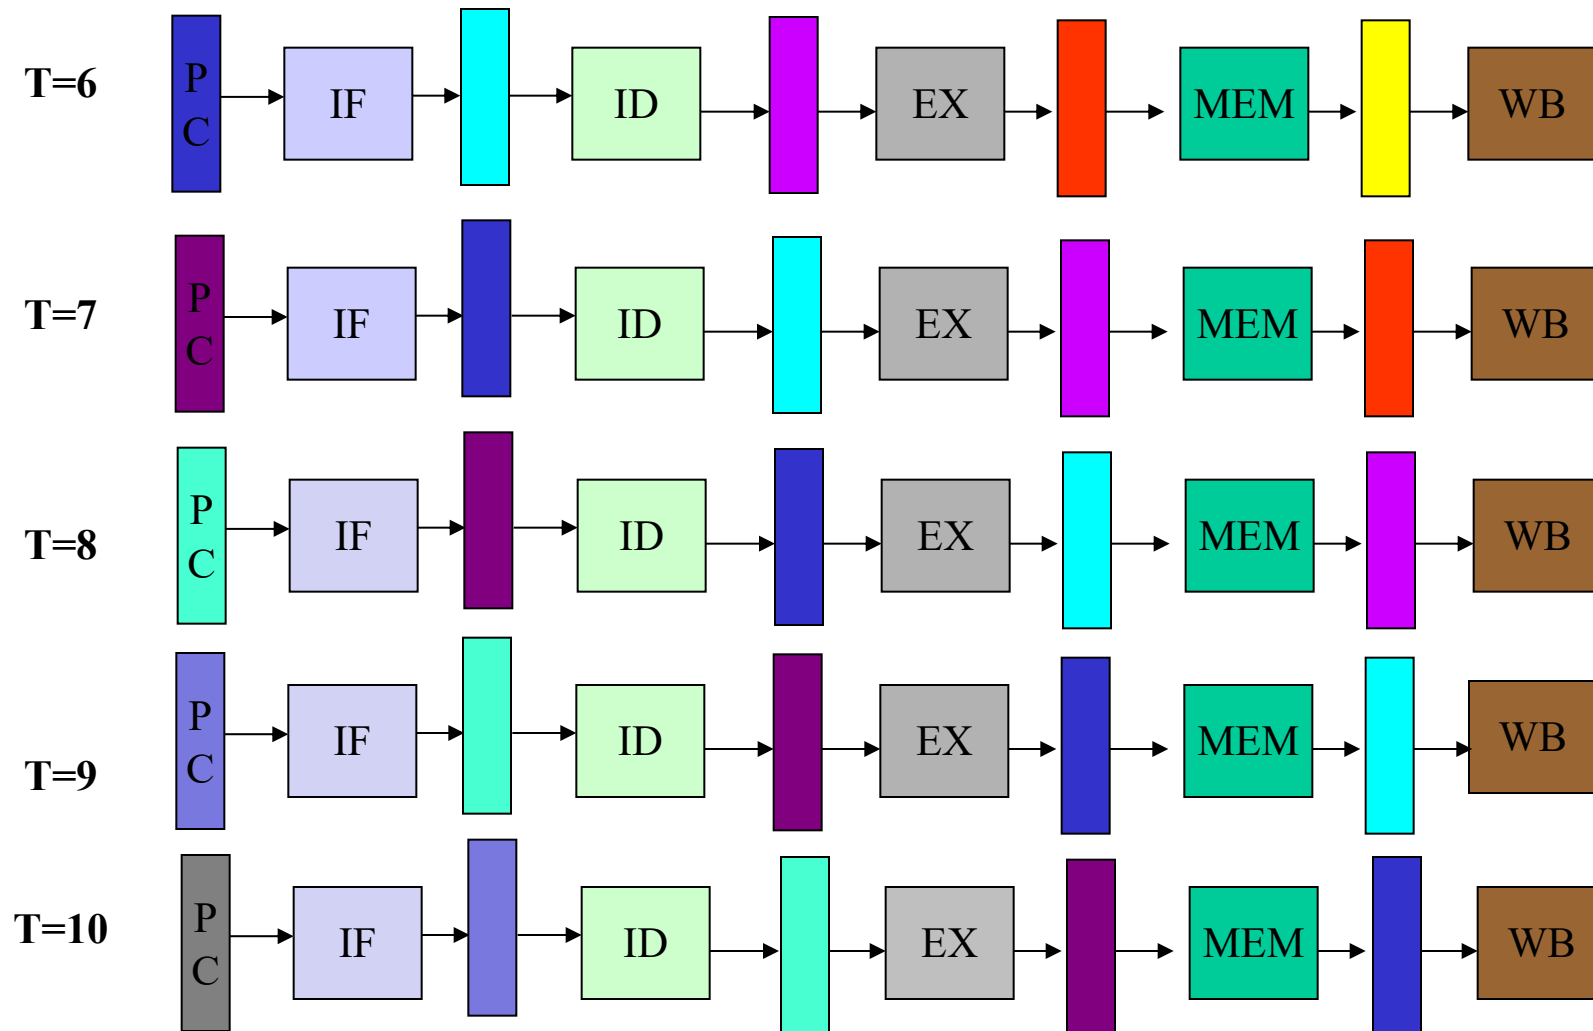
## 5-Stage Processor Pipeline



## 5-Stage Processor Pipeline



## 5-Stage Processor Pipeline



Instruction 6 will finish at cycle  $6 + 5 - 1 = 10$       Instruction  $n$  at cycle:  $n + 4$

## Ideal Pipeline Summary

$$\text{MIPS} = \text{IPC} \times \text{Clock Frequency (Hz)} \times 10^{-6}$$

$$\text{CPI} = (n+4 \text{ (cycles)}) / n \text{ (instructions)} = 1 + 4/n$$

Approximately  $\text{CPI} = 1$  for large  $n$  with no stalls

**Clock frequency** : Roughly 5 times that of base design

- Limited by the **delay through the pipeline stage**
  - Faster circuitry
  - Less circuitry with more stages
  - Will be ultimately limited by pipeline register delays
    - In practice useful functions require at least 10 gate delays
- Deep pipelines with very thin stages
  - **Clock rate** (frequency) **increase**
  - **Instruction throughput (MIPS)** **increases** in **ideal pipeline**
  - **Single instruction latency** **increase**
  - May cause greater number of stall cycles with data and control hazards



Blank Slide

## Pipeline Hazards

- What factors force deviations from this ideal pipeline?
  - Hazard
    - Condition that disrupts the orderly flow of instructions
    - Requires special attention by hardware and/or software to maintain program correctness and performance
1. Structural Hazards
    - Contention for hardware resources
  2. Data Hazards
    - Data dependencies between instructions
  3. Control Hazards
    - Disruptions caused by program control flow

## Structural Hazards

Contention for hardware resources

- **Memory Contention**

Instructions in IF and MEM stages may contend for memory

We will **assume** separate instruction and data memories to avoid conflict

- **ALU Contention**

Only 1 instruction doing an ALU operation at any time

**No contention** in this pipeline

- **Register File (RF) Contention**

- Instruction writes to RF while other instruction is reading RF

## Memory Contention

Instructions in IF and MEM stages may contend for memory

### Example

10% of instructions are LOADS and 5% are STORES.

A **single-ported unified** Instruction and Data Memory.

What is the **effective CPI**?

Each **LOAD and STORE** will **add 1 stall** cycle to the pipeline

$$\text{CPI} = 1.0 \text{ (nominal CPI)} + 10\% \times 1 + 5\% \times 1 = 1.15$$

We will **assume** separate instruction and data memories to avoid conflict

## Memory Contention

**Example 1:** 1% instructions encounter a Cache Miss in the IM.  
Miss Penalty 100 cycles. What is the effective CPI?

Each IM cache miss adds 100 cycle stall.  $CPI = 1.0 + 1\% \times 100 = 2.0$

**Example 2** 10% of instructions are LOADS and have a Data Miss Rate of 20% in DM. What is the effective CPI?

**Worst Case:** Miss in the IM and DM never occur together

A data miss adds 100 additional stall cycles i.e  $10\% \times 20\% \times 100 \text{ cycles} = 2 \text{ cycles}$

$$CPI = 2.0 + 2.0 = 4.0$$

**Best Case:** Assuming instruction and data misses can be served concurrently

All instructions misses (1%) overlap with a data miss (2%): All instruction miss stall times subsumed completely by data miss stalls.

Stall cycles is only for the data miss stalls: 2 cycles

$$CPI = 1.0 + 2.0 = 3.0$$